

M2-Thesis

Daniel Gagliardi

August 2026

Contents

Preface	3
1 Introduction	5
1.1 Initial Project	5
1.2 Existing work	5
1.3 Objects Investigated	6
1.4 Main Contribution	6
2 Background	7
3 Anisotropic Stability Bound	8
3.1 The deterministic transport mechanism	8
3.2 From the telescoping identity to the matrix bound	9
3.3 Recovering the one-sided Lipschitz estimate	10
3.4 A two-direction picture	10
3.5 What the estimate requires of the error	11
4 Level Sets	13
5 Ridges	16
6 Posterior Covariance Field	17
6.1 Three functionals of the spectrum	17
6.1.1 Effective dimension	17
6.1.2 Distance to projector structure	18
6.1.3 Ambiguity	19
6.2 MMSE-driven discretization	20
6.3 Examples	21
6.4 Summary	21

Preface

Throughout my undergraduate and graduate studies in mathematics I have been caught between two main interests: the applied world of machine learning, and the theoretical world of topology and geometry. While this is often viewed as an unusual mix to choose, I have always found it greatly enjoyable.

This culminated in me wanting to answer the question of what is happening to the “shape” of data throughout a machine learning process, and also what even is the “shape” of data.

The value in understanding this is twofold: first, to be able to make better models by leveraging our understanding of how they work, and second, to be able to understand how to interpret what these models are doing. The dance between naturally arising data and the model trying to fit without knowing how the data is generated is a curious one. My main motivation was to try to understand this.

One big idea is the notion of data lying on some low-dimensional space, which makes both learning and sampling from distributions supported on this lower-dimensional space easier. However this belief only truly simplifies the system in the small noise regime. All sorts of complex geometry can happen in the higher noise regimes. We do not, as we will see, have a nicely behaved interpolating family of manifolds.

Before doing the IASD M2 program, I had written a master’s dissertation on diffusion models under the manifold hypothesis, under the supervision of George Deligiannidis. During this I became deeply interested in dimension estimation of datasets. This at best vague and at worse terribly ill-posed question seemed to me one where an answer would tell us a lot about the natural geometry of data, and how diffusion models are able to make use of this. Almost every paper I found on the interplay between manifolds and diffusion models cited work from Eddie Aamari.

In November 2025 I reached out to Eddie Aamari to further discuss what the geometry of data could look like, and this led to this internship project. The initial goal was to understand the interplay between geometry affecting sampling error and how this could encourage simpler embeddings in a latent diffusion setting. However, given the complexity of work being done limited time was spent on investigating autoencoders so far.

A key challenge of this project has been striking the balance between geometry: where things are well defined, and diffusion of probability densities: where simply by adding noise we violently change the notions of geometry we would like to equip our density with. This tradeoff has led to looking at several objects, which while interesting in their own right have not been helpful in bounding sampling error. We do think the final part on posterior covariance fields [section 6](#) does provide new insight on where true complexity lies in the diffusion process.

Abstract

This is the abstract of the thesis.

Algebra is the offer made by the devil to the mathematician. The devil says: 'I will give you this powerful machine, it will answer any question you like. All you need to do is give me your soul: give up geometry and you will have this marvellous machine.'

— *Michael Atiyah*

1 Introduction

1.1 Initial Project

The progress in generative modelling over the last decade has been rapid. A key shift was the introduction of score based generative models (SGM), also known as diffusion models, which allow us to interpolate between a known distribution (usually a standard Gaussian), and a target distribution p_{data} , by training a network based on data $X \sim p_{\text{data}}$. By repeatedly adding noise to a target density, and training a denoising network, one can sample from p_{data} . Now, flow matching methods are the standard implementation of diffusion based generative models. These crucially can be viewed as a SGM.

While vast amounts of work has been done on improving the capabilities of diffusion models, the geometry of the entire diffusion process has been under explored.

Initially, this research internship aimed to investigate the geometry of both the autoencoder and denoiser in latent diffusion models. If one can show that the geometry of a target distribution influences the estimation of scores, discretisation error, or sampling stability, then the case can be made that when the autoencoder and diffusion model are trained jointly, that the autoencoder learns a more geometrically regular embedding.

We spent the internship so far almost entirely looking at the geometry of diffusion.

1.2 Existing work

There exists literature on how different geometric aspects affect the diffusion process, but these are mostly asymptotic results for highly regular target densities. [list some papers!!!]

The main challenge of the question “how does geometry affect diffusion” is what do we mean by geometry? There are some issues with existing methods which make use of dimension and reach/curvature of the support of p_{data} :

- Dimension:
 - Ill-posed to assign a dimension $d < D$ to a density which is absolutely continuous with respect to $\mathbb{R}^D$
 - Dimension of p_{data} may not be constant
 - Dimension of p_{data} may not be well defined everywhere, for example at intersections
 - Only considers dimension of target, now how effective dimension might change.
- Reach, curvature:
 - Can be $0, \infty$ at intersections or corners
 - Can vary wildly along the support of p_{data}

1.3 Objects Investigated

This gap in the literature is the primary of the project. As such, a new object other than dimension or reach or curvature had to be found. Throughout we investigated several before landing on the final and most promising object: the posterior covariance.

- Pullback metric of flow $G_{t,T}(x)$
- Level sets
- Ridges
- Posterior Covariance

[!!!make a table comparing them]

1.4 Main Contribution

We find that the spectrum of the posterior covariance captures geometric variation, and use functionals of the spectrum to understand geometric variation throughout the diffusion process. One of the things this allows us to do is pick a discretisation schedule to reduce the discretisation error in sampling.

2 Background

3 Anisotropic Stability Bound

We record a stability estimate for generator matching that keeps track of *which* directions of a drift error are amplified and which are damped. The scalar one-sided Lipschitz estimate propagates every error with a single worst-case rate; it is sharp only in an isotropic sense. The estimate below replaces that scalar rate by the derivative flow of the learnt dynamics, a matrix that sees the direction of the error. We isolate the mechanism in the pure-transport case first: throughout Sections 3.1–3.3 the diffusion coefficient is set to zero, and the additive-noise case is recovered verbatim by averaging the derivative flow over the Brownian trajectories (Remark 3.3).

3.1 The deterministic transport mechanism

Consider two time-dependent velocity fields $v_t, \hat{v}_t : \mathbb{R}^D \rightarrow \mathbb{R}^D$ for $0 \leq t \leq T$, smooth enough in space and integrable in time that their characteristic flows are globally defined and differentiable in the initial condition:

$$\frac{d}{ds} \Phi_{t,s}(x) = v_s(\Phi_{t,s}(x)), \quad \frac{d}{ds} \hat{\Phi}_{t,s}(x) = \hat{v}_s(\hat{\Phi}_{t,s}(x)), \quad \Phi_{t,t} = \hat{\Phi}_{t,t} = \text{Id}. \quad (1)$$

We write $\delta v_t := v_t - \hat{v}_t$ for the velocity error and think of v_t as the exact field and $\hat{v}_t$ as the learnt one.

The estimate is built entirely from the derivative of the *learnt* flow,

$$\hat{J}_{t,T}(x) := D_x \hat{\Phi}_{t,T}(x), \quad \frac{d}{ds} \hat{J}_{t,s}(x) = \nabla \hat{v}_s(\hat{\Phi}_{t,s}(x)) \hat{J}_{t,s}(x), \quad \hat{J}_{t,t}(x) = \text{Id}. \quad (2)$$

Thus $\hat{J}_{t,T}(x)h$ is the terminal displacement at time T produced by an infinitesimal perturbation h applied at (t, x) and transported by the learnt dynamics. The learnt flow acts as a magnifying glass on errors: an error committed at time t is not scaled by a worst-case constant but carried by $\hat{J}_{t,T}$ along its own direction. The following identity makes this exact.

Lemma 3.1 (Telescoping identity along characteristics). *For every initial point $x_0 \in \mathbb{R}^D$, with $X_t := \Phi_{0,t}(x_0)$ the exact characteristic,*

$$\Phi_{0,T}(x_0) - \hat{\Phi}_{0,T}(x_0) = \int_0^T \hat{J}_{t,T}(X_t) \delta v_t(X_t) dt. \quad (3)$$

Proof. Define $F(t) := \hat{\Phi}_{t,T}(X_t)$, which follows the exact dynamics up to time t and the learnt dynamics from t to T . Its endpoints are the fully learnt and fully exact terminal points, $F(0) = \hat{\Phi}_{0,T}(x_0)$ and $F(T) = \Phi_{0,T}(x_0)$. Differentiating in t ,

$$\frac{d}{dt} F(t) = \partial_t \hat{\Phi}_{t,T}(X_t) + D \hat{\Phi}_{t,T}(X_t) \dot{X}_t.$$

Starting the learnt flow an instant later removes one short learnt step, so $\partial_t \hat{\Phi}_{t,T}(x) = -D \hat{\Phi}_{t,T}(x) \hat{v}_t(x)$. Together with $\dot{X}_t = v_t(X_t)$ and $D \hat{\Phi}_{t,T}(X_t) = \hat{J}_{t,T}(X_t)$ this gives

$$\frac{d}{dt} F(t) = \hat{J}_{t,T}(X_t) (v_t(X_t) - \hat{v}_t(X_t)) = \hat{J}_{t,T}(X_t) \delta v_t(X_t).$$

Integrating $F(T) - F(0)$ from 0 to T yields (3). $\square$

Identity (3) is the heart of the estimate. The error made at each time enters the terminal gap through the derivative flow evaluated in the actual direction of that error, with no worst-case constant and no loss.

3.2 From the telescoping identity to the matrix bound

We turn the pathwise identity into a bound on the terminal laws by coupling. Let $X_0 \sim \rho_0$ and transport it by both flows, $X_t := \Phi_{0,t}(X_0)$ and $\hat{X}_t := \hat{\Phi}_{0,t}(X_0)$, with terminal laws $\rho_T := \text{Law}(X_T)$ and $\hat{\rho}_T := \text{Law}(\hat{X}_T)$. Running both flows from the same X_0 is a coupling of ρ_T and $\hat{\rho}_T$, so $W_1(\rho_T, \hat{\rho}_T) \leq \mathbb{E}\|X_T - \hat{X}_T\|$. Applying Lemma 3.1 pathwise at $x_0 = X_0$ and the triangle inequality,

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \mathbb{E} \|\hat{J}_{t,T}(X_t) \delta v_t(X_t)\| dt = \int_0^T \int_{\mathbb{R}^D} \|\hat{J}_{t,T}(x) \delta v_t(x)\| \rho_t(dx) dt, \quad (4)$$

where the last step uses $X_t \sim \rho_t$. It remains only to name the quadratic form hidden in the norm.

Theorem 3.2 (Anisotropic W_1 stability). *Define the pullback amplification metric*

$$G_{t,T}(x) := \hat{J}_{t,T}(x)^\top \hat{J}_{t,T}(x), \quad \text{so that} \quad \|\hat{J}_{t,T}(x)h\|^2 = \langle h, G_{t,T}(x)h \rangle. \quad (5)$$

Then

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \int_{\mathbb{R}^D} \langle \delta v_t(x), G_{t,T}(x) \delta v_t(x) \rangle^{1/2} \rho_t(dx) dt. \quad (6)$$

Proof. Substitute the identity in (5) into (4). $\square$

The matrix $G_{t,T}(x)$ is the Euclidean metric at the terminal time pulled back to the tangent space at (t, x) by the learnt flow: it lives at time t but measures squared lengths after transport to T . Its eigenvectors are the input directions at (t, x) and its eigenvalues are the squared stretching factors of those directions. In the deterministic case (5) is an exact identity, so no Cauchy–Schwarz slack is incurred in passing from the derivative flow to the quadratic form. The bound (6) is anisotropic precisely because it evaluates $G_{t,T}$ on the actual error direction $\delta v_t(x)$ rather than on the worst possible direction.

Remark 3.3 (Additive diffusion). When both dynamics carry the same additive diffusion coefficient, the Brownian term has no derivative in the initial condition, so the learnt derivative flow $\hat{J}_{t,s}$ still solves the linear equation (2) along each trajectory. The metric becomes the trajectory average $G_{t,T}(x) = \mathbb{E}[\hat{J}_{t,T}(x)^\top \hat{J}_{t,T}(x)]$, and Jensen’s inequality turns the exact identity in (5) into the upper bound $\mathbb{E}\|\hat{J}_{t,T}(x)\delta v_t(x)\| \leq \langle \delta v_t, G_{t,T}\delta v_t \rangle^{1/2}$, so (6) holds unchanged. The deterministic case treated here is the load-bearing one.

3.3 Recovering the one-sided Lipschitz estimate

The classical scalar estimate is the isotropic specialization of Theorem 3.2, obtained by discarding the direction of the error. Let $S_t(x) := \frac{1}{2}(\nabla \hat{v}_t(x) + \nabla \hat{v}_t(x)^\top)$ be the symmetric part of the learnt Jacobian, and suppose the learnt field is one-sided Lipschitz with rate ℓ_t , meaning

$$\lambda_{\max}(S_t(x)) \leq \ell_t \quad \text{for all } x. \quad (7)$$

The condition constrains expansion only: it bounds the largest rate at which two trajectories separate, and places no bound on contraction or on the rotational part of $\nabla \hat{v}_t$.

Corollary 3.4 (Scalar one-sided Lipschitz bound). *Under (7),*

$$G_{t,T}(x) \preceq \exp\left(2 \int_t^T \ell_s \, ds\right) \text{Id}, \quad (8)$$

and consequently

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \exp\left(\int_t^T \ell_s \, ds\right) \|\delta v_t\|_{L^1(\rho_t)} \, dt. \quad (9)$$

Proof. Fix a direction h and set $z_s := \hat{J}_{t,s}(x)h$, which solves $\dot{z}_s = \nabla \hat{v}_s(\hat{\Phi}_{t,s}(x)) z_s$ by (2). Only the symmetric part contributes to the growth of the norm, so

$$\frac{d}{ds} \|z_s\|^2 = 2 \langle z_s, S_s(\hat{\Phi}_{t,s}(x)) z_s \rangle \leq 2\ell_s \|z_s\|^2.$$

Grönwall's lemma gives $\|\hat{J}_{t,T}(x)h\| \leq \exp(\int_t^T \ell_s \, ds) \|h\|$, which is (8). Inserting (8) into (6) and using $\langle \delta v_t, G_{t,T} \delta v_t \rangle^{1/2} \leq \exp(\int_t^T \ell_s \, ds) \|\delta v_t\|$ yields (9). $\square$

Corollary 3.4 exhibits the scalar estimate as the Loewner domination (8): the matrix $G_{t,T}$ is replaced by its largest eigenvalue, uniformly over directions. The bound is correct, but it forgets the eigenstructure of $G_{t,T}$, and in particular treats a strongly contracted error direction as though it were the most expansive one. The next subsection makes this gap concrete in a two-direction model, where a normal error is carried by a contraction factor that the scalar constant cannot see.

3.4 A two-direction picture

The gap between the matrix and scalar bounds is already visible in the simplest anisotropic model. Fix an orthogonal splitting $\mathbb{R}^D = E_{\parallel} \oplus E_{\perp}$ into a tangent-like block and a normal-like block, write $x = (x_{\parallel}, x_{\perp})$ accordingly, and take the learnt velocity to be linear, block-diagonal, expansive on $E_{\parallel}$ and contractive on $E_{\perp}$:

$$\hat{v}_t(x_{\parallel}, x_{\perp}) = (\ell_t x_{\parallel}, -\gamma_t x_{\perp}), \quad \gamma_t \geq 0. \quad (10)$$

Its Jacobian $\nabla \hat{v}_t = \text{diag}(\ell_t I_{\parallel}, -\gamma_t I_{\perp})$ is constant in space, so the derivative flow (2) integrates blockwise and the pullback metric (5) is diagonal:

$$\hat{J}_{t,T} = \begin{pmatrix} e^{\int_t^T \ell_s \, ds} I_{\parallel} & 0 \\ 0 & e^{-\int_t^T \gamma_s \, ds} I_{\perp} \end{pmatrix}, \quad G_{t,T} = \begin{pmatrix} e^{2 \int_t^T \ell_s \, ds} I_{\parallel} & 0 \\ 0 & e^{-2 \int_t^T \gamma_s \, ds} I_{\perp} \end{pmatrix}. \quad (11)$$

Splitting the error as $\delta v_t = \delta v_t^{\parallel} + \delta v_t^{\perp}$ with $\delta v_t^{\parallel} \in E_{\parallel}$ and $\delta v_t^{\perp} \in E_{\perp}$, the integrand of the matrix bound (6) separates,

$$\langle \delta v_t, G_{t,T} \delta v_t \rangle^{1/2} = \left(e^{2 \int_t^T \ell_s \, ds} \|\delta v_t^{\parallel}\|^2 + e^{-2 \int_t^T \gamma_s \, ds} \|\delta v_t^{\perp}\|^2 \right)^{1/2}. \quad (12)$$

A tangent error is carried to the terminal time by the expansion factor $e^{\int_t^T \ell}$; a normal error is carried by the contraction factor $e^{-\int_t^T \gamma}$ and is therefore damped.

The scalar estimate cannot distinguish the two. Its rate is the largest eigenvalue of the symmetric Jacobian,

$$\lambda_{\max}(S_t) = \max\{\ell_t, -\gamma_t\} = \ell_t \quad (\ell_t \geq 0), \quad (13)$$

so Corollary 3.4 propagates *both* error components with the tangent expansion factor $e^{\int_t^T \ell}$. It replaces the normal contraction $e^{-\int_t^T \gamma}$ by expansion, overestimating the contribution of a normal error by the factor $e^{\int_t^T (\ell_s + \gamma_s) \, ds}$. When the normal block is strongly contracting—as it is near a terminal manifold, where $\gamma_t \sim (1-t)^{-1}$ and $e^{-\int_t^T \gamma}$ decays polynomially in $1-t$ —this overestimate is the entire difference between a persistent error and a damped one. This is where the anisotropy pays.

3.5 What the estimate requires of the error

The two bounds depend on the error through different amounts of information, and this raises a natural objection. The scalar bound (9) uses only the magnitude $\|\delta v_t\|_{L^1(\rho_t)}$. The matrix bound (6) uses the quadratic form $\langle \delta v_t, G_{t,T} \delta v_t \rangle$, that is, the *direction* of δv_t relative to the eigenframe of $G_{t,T}$. Since the exact field v_t is unknown, the error $\delta v_t = v_t - \hat{v}_t$ is never observed, and its direction still less so. It is not immediate that the matrix bound is any more usable than the scalar one.

The resolution is that neither bound is an estimator: both are a priori certificates. The scalar bound is valid because it implicitly places δv_t along the most expansive eigenvector of $G_{t,T}$, which is exactly the source of its looseness. The matrix bound is a family of certificates indexed by hypotheses on where the error lies, and it becomes informative as soon as one can argue, on structural grounds, that the error concentrates in contracted directions. The direction is supplied not by measuring δv_t but by the geometry of the *target*.

Near a smooth embedded manifold this structure is explicit. The exact velocity v_t , which is fixed by the score of the noised target, admits to leading order a canonical decomposition: a rigid normal restoring field of model-independent

functional form, a tangential component set by the intrinsic density on the manifold, and a moving ridge that locates the normal zero of the field. Any learnt error is then graded in the *same* frame into a normal-amplitude component, a tangential-density component, and a ridge-displacement component. The matrix bound assigns each its own weight through $G_{t,T}$: normal-amplitude errors are damped by the normal contraction factor of [subsection 3.4](#), tangential-density errors are carried by the tangential expansion factor and persist, and ridge-displacement errors are normal but enter with the singular weight σ_t^{-2} , since a normal displacement of the ridge by Δ changes the leading normal field by Δ/σ_t^2 . The estimate never requires δv_t to be known; it requires the target to furnish a fixed anisotropic frame in which any error is automatically decomposed. We construct this frame in [section 6](#).

One assumption must be stated now, because it is used throughout. The projectors defining the frame— P_t onto the tangent block and $Q_t = I - P_t$ onto the normal block—must be a single field, determined by the target and shared across all learnt fields. Re-deriving the frame from each model’s own error would reintroduce model dependence into the weights and render comparisons between models meaningless. We therefore fix P_t once, from the target geometry, and read every error in that frame; concretely, P_t is built from the spectrum of the posterior-covariance operator A_t in [section 6](#).

4 Level Sets

as we smooth data which lies on or is concentrated around a manifold two main affects occur. Smoothing in the normal direction and smoothing and tangent direction. Smoothing in the tangent direction makes a density more uniform along a manifold. Therefore our level sets should be better behaved.

Let $p^\star : \mathbb{R}^D \rightarrow \mathbb{R}_+$ denote the initial density, and let p_σ denote its Gaussian-smoothed version,

$$p_\sigma(x) = (G_\sigma * p^\star)(x), \quad G_\sigma(x) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right). \quad (14)$$

For a regular value c , the corresponding level set is

$$\mathcal{M}_c^\sigma := \{x \in \mathbb{R}^D : p_\sigma(x) = c\}, \quad \nabla p_\sigma(x) \neq 0 \quad \text{for } x \in \mathcal{M}_c^\sigma. \quad (15)$$

The unit normal to $\mathcal{M}_c^\sigma$ is

$$n(x) = \frac{\nabla p_\sigma(x)}{\|\nabla p_\sigma(x)\|}. \quad (16)$$

For a tangent direction $v \in T_x \mathcal{M}_c^\sigma$, the derivative of the normal is

$$D_v n = \frac{1}{\|\nabla p_\sigma\|} (I - nn^\top) D^2 p_\sigma v = \frac{1}{\|\nabla p_\sigma\|} P_T D^2 p_\sigma v, \quad (17)$$

where

$$P_T := I - nn^\top \quad (18)$$

is the orthogonal projection onto the tangent space $T_x \mathcal{M}_c^\sigma$. Thus, the variation of the normal field—and hence the local curvature of the level set—is determined by the Hessian of the density.

Gaussian smoothing regularizes not only the density itself, but all of its spatial derivatives. In particular, for any $k \geq 0$,

$$D^k p_\sigma = D^k (G_\sigma * p^\star) = (D^k G_\sigma) * p^\star. \quad (19)$$

Hence increasing σ smooths the density and simultaneously suppresses fine-scale variation in its gradient, Hessian, and higher-order derivatives.

This smoothing has a particularly simple interpretation in the Fourier domain. Writing $\hat{p}_\sigma$ for the Fourier transform,

$$\hat{p}_\sigma(\xi) = \hat{G}_\sigma(\xi) \hat{p}^\star(\xi) = \exp\left(-\frac{\sigma^2 \|\xi\|^2}{2}\right) \hat{p}^\star(\xi). \quad (20)$$

Thus a Fourier mode of spatial frequency $\|\xi\|$ is attenuated by $\exp(-\sigma^2 \|\xi\|^2/2)$: high-frequency components decay quadratically faster in frequency. Consequently, Gaussian smoothing preferentially removes the fine-scale spatial variation of $p^\star$.

The effect on the geometry of a level set follows directly from (17). Its shape operator is

$$S := -P_T Dn P_T = -\frac{1}{\|\nabla p_\sigma\|} P_T D^2 p_\sigma P_T, \quad (21)$$

up to the choice of orientation for the normal. The eigenvalues of S are the principal curvatures of $\mathcal{M}_c^\sigma$. Therefore, Gaussian smoothing affects the shape of the level set through two quantities: the Hessian $D^2 p_\sigma$, which controls variation of the normal, and $\|\nabla p_\sigma\|$, which sets the scale relating density variation to geometric variation. In particular, since Gaussian smoothing suppresses high-frequency components of $D^2 p_\sigma$, it suppresses the fine-scale variation of the shape operator and hence removes small-scale geometric structure from the level set.

It is important to distinguish this suppression of fine-scale variation from a pointwise reduction in curvature. In general, Gaussian smoothing does not imply that the principal curvatures of a fixed level set decrease monotonically with σ . Indeed, the shape operator

$$S_\sigma(x) = -\frac{1}{\|\nabla p_\sigma(x)\|} P_T D^2 p_\sigma(x) P_T, \quad x \in \mathcal{M}_c^\sigma, \quad (22)$$

depends on both $D^2 p_\sigma$ and ∇p_σ , and is evaluated on the moving level set $\mathcal{M}_c^\sigma$. Although Gaussian smoothing suppresses high-frequency components of $D^2 p_\sigma$, it simultaneously changes ∇p_σ and the location of the level set itself. Hence there is no general pointwise monotonicity of S_σ or of its eigenvalues.

Moreover, the topology of $\mathcal{M}_c^\sigma$ can change with σ . Such a change occurs when the level c passes through a critical point, i.e., when

$$p_\sigma(x) = c, \quad \nabla p_\sigma(x) = 0. \quad (23)$$

For example, two components of a superlevel set may merge through a neck as σ increases. Thus, while smoothing removes fine-scale spatial variation from the density and its derivatives, it should not be interpreted as a monotone decrease of the curvature of every individual level set. The rigorous statement is instead that Gaussian smoothing preferentially suppresses spatial frequencies at scales smaller than σ , and therefore suppresses fine-scale geometric variation encoded by the derivatives of the density.

More generally, adding Gaussian noise monotonically suppresses all L^2 Sobolev energies of the smoothed density. Let

$$p_\sigma = G_\sigma * p^*, \quad G_\sigma(x) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right), \quad (24)$$

where D is the ambient dimension. Since $\partial_{\sigma^2} p_\sigma = \frac{1}{2} \Delta p_\sigma$, we have, for every integer $k \geq 0$,

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int_{\mathbb{R}^D} |D^k p_\sigma(x)|^2 dx = -\frac{1}{2} \int_{\mathbb{R}^D} |\nabla D^k p_\sigma(x)|^2 dx \leq 0. \quad (25)$$

Thus, increasing the noise level σ monotonically decreases the L^2 norms of all spatial derivatives of p_σ . In particular,

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int |\nabla p_\sigma|^2 dx = -\frac{1}{2} \int |\Delta p_\sigma|^2 dx \leq 0, \quad (26)$$

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int |D^2 p_\sigma|^2 dx = -\frac{1}{2} \int |\nabla \Delta p_\sigma|^2 dx \leq 0. \quad (27)$$

Equivalently, in the Fourier domain,

$$\frac{1}{2} \int_{\mathbb{R}^D} |D^k p_\sigma(x)|^2 dx \propto \frac{1}{2} \int_{\mathbb{R}^D} \|\xi\|^{2k} e^{-\sigma^2 \|\xi\|^2} |\widehat{p^\star}(\xi)|^2 d\xi. \quad (28)$$

Hence increasing σ progressively removes increasingly fine-scale structure, with spatial frequency $\|\xi\|$ suppressed according to $e^{-\sigma^2 \|\xi\|^2}$.

This relates to the conventional wisdom where we need more refined time steps as we approach closer to our final density to capture the higher frequency structure of density. This is particularly apparent for image and audio data. JPEG compression takes advantage of writing the image as a Fourier decomposition, and in doing so discard the high frequency terms which are not noticeable to the human eye.

5 Ridges

An at first unintuitive aspect of the manifold hypothesis is that the lower dimensional space around which data concentrate is not necessarily, or in fact usually, the initial manifold. This is because as the amount of noise that we add increases the effective radius with which we convolve grows and includes more of our initial manifold. We see a noise density concentrating around a different space. This begs the question “what is this new space and how does it behave and evolve?”

This becomes immediately obvious when considering smoothing a pair of Gaussians: the modes of the sum is not the union of the modes.

below as an example of what smoothing a density on a sine wave can do.

With this intuition it’s fairly easy to rediscover the definition of a rich. It’s where the normal component of our changing density is zero. But a priori we do not know which components are normal.

6 Posterior Covariance Field

Here we introduce the ultimate object we analyse: posterior covariance fields.

Definition 6.1. Let $Y_t = X + \sqrt{t}Z$, $m_t(y) := \mathbb{E}[X \mid Y_t = y]$, and

$$A_t(y) := \frac{1}{t} \text{Cov}(X \mid Y_t = y).$$

We call the weighted object $\mathfrak{F}_t := \text{Law}(Y_t, m_t(Y_t), A_t(Y_t))$ the *posterior covariance field*.

Note: while we are also including the posterior mean $m_t(y)$ in our definition, we keep the name for simplicity. The posterior mean m_t and by extension the score s_t is the object whose regularity we want to understand, so its inclusion is natural.

There are a few reasons why it makes sense to look at this object. First, it is not plagued by the issues we encountered in the analysis of ridges or level sets in [section 5](#) and [section 4](#). These objects suffered from being too geometric for our use case, where we care about noise scales much greater than 0. We need a more analytic quantity, and the posterior covariance field is defined for all noise scales $t > 0$. Second, the posterior covariance captures the local geometry of the score. The posterior covariance corresponds to the Jacobian of the score. This link is given in [section 2](#). Here we normalise the posterior covariance by t to account for the scale of noise added. Third, here we deal with expectations, rather than L^∞ -type quantities like the reach τ_M . If there exists a region of the density which is geometrically complex and as a consequence has poor regularity, but which carries almost no mass, we should not let it dramatically change our understanding of the behaviour of the score.

6.1 Three functionals of the spectrum

All three quantities below are moments of the same object: the spectral measure of A_t , averaged over $y \sim p_t$. Write $a_1(y), \dots, a_D(y)$ for the eigenvalues of $A_t(y)$ and let $\mu_t := \text{Law}(a_I(Y_t))$ with I uniform on $\{1, \dots, D\}$, so that $\mathbb{E} \sum_i f(a_i) = D \int f d\mu_t$ for any f .

6.1.1 Effective dimension

Definition 6.2.

$$d_{\text{eff}}(t) := \mathbb{E}[\text{tr } A_t(Y_t)] = \mathbb{E} \sum_i a_i. \quad (29)$$

Proposition 6.3 (Exact MMSE identity). *Let $\text{mmse}(t) := \mathbb{E}|X - m_t(Y_t)|^2$. Then*

$$d_{\text{eff}}(t) = \frac{\text{mmse}(t)}{t}. \quad (30)$$

Proof. $\text{tr Cov}(X \mid Y_t) = \mathbb{E}[|X - m_t(Y_t)|^2 \mid Y_t]$; take expectations and divide by t . $\square$

Identity (30) is exact, not asymptotic, and has two consequences of practical importance. First, d_{eff} is dimensionless: it is residual posterior uncertainty measured in units of injected noise variance. Second, $\text{mmse}(t)$ is precisely the population value of the ε -prediction (or x_0 -prediction) denoising loss at noise level t , so d_{eff} requires no computation beyond the training objective already being evaluated.

Link to dimension. An orthogonal projector of rank k has spectrum $\{1^{(k)}, 0^{(D-k)}\}$ and trace k . In the ideal manifold limit $A_t \approx P_{T_x \mathcal{M}}$ and hence $d_{\text{eff}}(t) \approx d$. Equivalently, via (30): a d -dimensional manifold leaves $O(t)$ posterior uncertainty in each of d tangent directions and $o(t)$ in the normal directions, so $\text{mmse} \approx d \times t$. Each eigenvalue may be read as a fractional dimension: $a_i = 0.4$ contributes 0.4 effective directions.

Failure mode. d_{eff} is a first moment and therefore cannot distinguish $\text{diag}(1, 1, 1, 0, 0, 0)$ from $\text{diag}(\frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2})$: both have trace 3. It also over-counts when $a_i > 1$, in which case the “fractional dimension” reading breaks down entirely. The next two functionals repair exactly these two deficiencies.

6.1.2 Distance to projector structure

Definition 6.4.

$$\Delta_{\text{proj}}(t) := \mathbb{E} \sum_i \min\{a_i^2, (a_i - 1)^2\}. \quad (31)$$

The summand is the squared distance from a_i to the set $\{0, 1\}$, and the two branches exchange at $a_i = \frac{1}{2}$:

$$\min\{a^2, (a - 1)^2\} = \begin{cases} a^2, & a \leq \frac{1}{2} \quad (\text{read as normal-like}), \\ (a - 1)^2, & a \geq \frac{1}{2} \quad (\text{read as tangent-like}). \end{cases} \quad (32)$$

Proposition 6.5. Let $\mathcal{P} = \{P \in \mathbb{R}^{D \times D} : P = P^\top = P^2\}$ be the set of orthogonal projectors and let A be symmetric with eigenvalues a_i . Then

$$\min_{P \in \mathcal{P}} \|A - P\|_F^2 = \sum_i \min\{a_i^2, (a_i - 1)^2\}, \quad \text{hence} \quad \Delta_{\text{proj}}(t) = \mathbb{E}[\text{dist}_F^2(A_t(Y_t), \mathcal{P})]. \quad (33)$$

Proof. $\mathcal{P}$ is invariant under conjugation by orthogonal matrices, and for symmetric A the Frobenius distance to a conjugation-invariant set is minimized by a P sharing the eigenvectors of A (a consequence of the Hoffman–Wielandt inequality). It then remains to round each a_i to the nearer of 0 or 1. $\square$

Thus Δ_{proj} answers: *how far, on average, is the soft posterior geometry from being a genuine tangent-space projector?* It vanishes exactly when $A_t(Y_t)$ is a projector almost surely.

Remark 6.6 (Interpretation as a resolution probe). For data concentrated near a manifold but with normal thickness δ , the normal eigenvalues behave as $a_n \approx \delta^2/(\delta^2 + t)$, which passes through $\frac{1}{2}$ at $t \approx \delta^2$ and contributes $\frac{1}{4}$ to Δ_{proj} there. Consequently $\Delta_{\text{proj}}(t)$ peaks at the scale at which the manifold hypothesis stops being exactly true, and its peak location is a readout of that scale. Δ_{proj} is best regarded as a scale-resolved test of manifold-likeness rather than a single verdict.

6.1.3 Ambiguity

Definition 6.7. With $(x)_+ := \max\{x, 0\}$,

$$\mathcal{A}(t) := \mathbb{E} \sum_i (a_i - 1)_+^2. \quad (34)$$

The motivation is dynamical rather than geometric. $a_i > 1$ occurs if and only if the reverse probability-flow ODE is locally expansive along v_i . The threshold is the separatrix at which posterior variance exceeds the injected noise variance:

$$a_i = \frac{\text{posterior variance along } v_i}{\text{noise variance } t} \quad \begin{cases} < 1 : & \text{denoising, contraction,} \\ = 1 : & \text{neutral,} \\ > 1 : & \text{amplification, expansion.} \end{cases} \quad (35)$$

The three design choices are then: threshold at 1 (only expansion is penalized), measure the *excess* $a_i - 1$ (the neutral point is 1, not 0), and square it (making the functional comparable to a squared Frobenius norm and penalizing strong expansion superlinearly).

Link to ambiguity. Suppose the posterior at y is approximately $\frac{1}{2}\mathcal{N}(\mu_1, \Sigma) + \frac{1}{2}\mathcal{N}(\mu_2, \Sigma)$, corresponding to two competing clean configurations. Then

$$\text{Cov}(X \mid Y_t = y) = \Sigma + \frac{1}{4}(\mu_1 - \mu_2)(\mu_1 - \mu_2)^\top, \quad (36)$$

and the between-mode term can exceed tI along $\mu_1 - \mu_2$, giving $a_i > 1$. So $a_i > 1$ signals that the noisy observation fails to identify a single locally stable latent configuration.

Remark 6.8 (Scope of the interpretation). Strictly, $a_i > 1$ is a diagnostic that posterior variance exceeds the noise scale. Multimodality is the typical, but not the only, mechanism: heavy-tailed unimodal p_0 (e.g. Laplace, Student- t) also produces $a_i > 1$. “Ambiguity” is the geometric reading of the phenomenon, not a theorem.

Remark 6.9 (Mean versus tail). $\mathcal{A}$ is an average over $y \sim p_t$ and can be small while $\sup_y a_{\text{max}}(y)$ is large, if the expansive region carries little mass. In practice one should report both $\mathcal{A}(t)$ and a high quantile of $(a_{\text{max}} - 1)_+$, since sampler failure is driven by trajectories that enter the expansive region, not by its measure.

6.2 MMSE-driven discretization

We started by looking at (Potapchik et al., 2024).

We consider the additive Gaussian perturbation

$$X_t = X_0 + \sqrt{t} Z, \quad Z \sim \mathcal{N}(0, I_D), \quad (37)$$

so that, in the notation of the diffusion model,

$$c_t = 1, \quad \sigma_t^2 = t. \quad (38)$$

Let

$$m_t(x) := \mathbb{E}[X_0 \mid X_t = x] \quad (39)$$

denote the posterior mean and define the minimum mean-squared error

$$\text{MMSE}(t) := \mathbb{E} [\|X_0 - m_t(X_t)\|^2]. \quad (40)$$

Following the effective-dimension viewpoint, we write

$$d_{\text{eff}}(t) := \frac{\text{MMSE}(t)}{t}, \quad (41)$$

and hence

$$\text{MMSE}(t) = t d_{\text{eff}}(t). \quad (42)$$

Discretization error. For additive Gaussian noise, Tweedie's formula gives

$$s_t(x) = \frac{m_t(x) - x}{t}. \quad (43)$$

The same martingale argument used to control the score variation in the discretized reverse process then yields, for two noise levels $t < t'$,

$$\begin{aligned} & \mathbb{E} \|s_t(X_t \mid t', X_{t'}) - s_t(X_t)\|^2 \\ &= \frac{1}{t^2} (\text{MMSE}(t') - \text{MMSE}(t)), \end{aligned} \quad (44)$$

with the sign understood according to the direction in which the diffusion trajectory is traversed. Thus, for a sufficiently fine discretization with step size

$$\Delta t_k := t_{k+1} - t_k, \quad (45)$$

the contribution of the k -th interval is, up to higher-order terms,

$$\mathcal{E}_k \asymp \frac{|\text{MMSE}'(t_k)|}{t_k^2} (\Delta t_k)^2. \quad (46)$$

Consequently, the total discretization error is approximated by

$$\mathcal{E}_{\text{disc}} \asymp \sum_{k=0}^{K-1} a(t_k) (\Delta t_k)^2, \quad a(t) := \frac{|\text{MMSE}'(t)|}{t^2}. \quad (47)$$

Optimal allocation of time steps. We now choose the discretization points to minimize (47) subject to a fixed number K of steps. Introduce the local step density

$$n(t) := \frac{1}{\Delta t(t)}. \quad (48)$$

In the continuum limit,

$$K \approx \int n(t) dt, \quad (49)$$

while

$$\begin{aligned} \mathcal{E}_{\text{disc}} &\approx \int a(t) \Delta t(t) dt \\ &= \int \frac{a(t)}{n(t)} dt. \end{aligned} \quad (50)$$

Minimizing this functional subject to $\int n(t) dt = K$ gives

$$n(t) \propto \sqrt{a(t)}, \quad (51)$$

and therefore

$$\Delta t(t) \propto \frac{1}{\sqrt{a(t)}} = \frac{t}{\sqrt{|\text{MMSE}'(t)|}}. \quad (52)$$

Finally, using (42),

$$\text{MMSE}'(t) = \frac{d}{dt} [t d_{\text{eff}}(t)] = d_{\text{eff}}(t) + t d'_{\text{eff}}(t), \quad (53)$$

so the resulting MMSE-adaptive discretization is

$$\boxed{\Delta t(t) \propto \frac{t}{\sqrt{|d_{\text{eff}}(t) + t d'_{\text{eff}}(t)|}}}. \quad (54)$$

This prescription allocates more discretization points to regions in which the posterior uncertainty, measured by the MMSE, changes rapidly, while accounting for the t^{-2} amplification of score errors at small noise scales.

6.3 Examples

6.4 Summary

References

Peter Potaptchik, Iskander Azangulov, and George Deligiannidis. Linear convergence of diffusion models under the manifold hypothesis. *arXiv preprint arXiv:2410.09046*, 2024.