

M2-Thesis

Daniel Gagliardi

August 2026

Contents

Preface	3
1 Introduction	5
1.1 Initial Project	5
1.2 Existing work	5
1.3 Objects Investigated	6
1.4 Main Contribution	6
2 Background	7
2.1 The Manifold Hypothesis	7
2.2 Flow Matching as a Score-Based Generative Model	8
2.3 Tweedie's 1 st and 2 nd Order Formulae	9
2.3.1 First Derivative (Score Function)	9
2.3.2 Second Derivative (Hessian Matrix/Score Jacobian)	10
2.3.3 Posterior Covariance Relation	10
3 Anisotropic Stability Bound	11
3.1 The deterministic transport mechanism	11
3.2 From the telescoping identity to the matrix bound	12
3.3 Recovering the one-sided Lipschitz estimate	12
3.4 A two-direction picture	13
3.5 What the estimate requires of the error	14
4 Level Sets	15
5 Ridges	18
5.1 Motivation	18
5.2 The score concentrates on the ridge	19
5.3 Failings of the Ridge	20
6 Posterior Covariance Field	22
6.1 Three functionals of the spectrum	22
6.1.1 Effective dimension	22
6.1.2 Distance to projector	23
6.1.3 Ambiguity	24
6.2 Evolution law of MMSE	24
6.3 MMSE-driven discretisation	25
6.4 Interpretation	26

Preface

Throughout my undergraduate and graduate studies in mathematics I have been caught between two main interests: the applied world of machine learning, and the theoretical world of topology and geometry. While this is often viewed as an unusual mix to choose, I have always found it greatly enjoyable.

This resulted in me wanting to answer the question of what is happening to the “shape” of data throughout a machine learning process, and also what even is the “shape” of data.

The value in understanding this is twofold: firstly, to be able to make better models by leveraging our understanding of how they work; and second, to be able to understand how to interpret what these models are doing. The dance between naturally arising data and the model trying to fit without knowing how the data is generated is fascinating. My main motivation was to try to understand this.

One big idea is the notion of data lying on some low-dimensional space, which makes both learning and sampling from distributions supported on this lower-dimensional space easier. However this belief only truly simplifies the system in the small noise regime. All sorts of complex geometry can happen in the higher noise regimes. We do not, as we will see, have a nicely behaved interpolating family of manifolds.

Before doing the IASD M2 program, I had written a master’s dissertation on diffusion models under the manifold hypothesis, under the supervision of George Deligiannidis. During this I became deeply interested in dimension estimation of datasets. This at best vague and at worst terribly ill-posed question seemed to me one where an answer would tell us a lot about the natural geometry of data, and how diffusion models are able to make use of this. Almost every paper I found on the interplay between manifolds and diffusion models cited work from Eddie Aamari.

In November 2025 I reached out to Eddie Aamari to further discuss what the geometry of data could look like, and this led to this internship project. The initial goal was to understand the interplay between geometry affecting sampling error and how this could encourage simpler embeddings in a latent diffusion setting. However, given the complexity of work being done limited time was spent on investigating autoencoders so far.

A key challenge of this project has been striking the balance between geometry: where things are well defined, and diffusion of probability densities: where simply by adding noise we violently change the notions of geometry we would like to equip our density with. This tradeoff has led to looking at several objects, which while interesting in their own right have not been helpful in bounding sampling error. I do believe that the final part on posterior covariance fields [section 6](#) does provide a new insight on where true complexity lies in the diffusion process.

Abstract

Diffusion based generative models have become ubiquitous over recent years. Their success is often attributed to the geometric structure of natural data. This is commonly referred to as the manifold hypothesis. In this report we summarise our main work so far on understanding the geometry of diffusion. We examine an anisotropic stability bound, level sets, ridges, and functionals of the posterior covariance. Ultimately, we settle on the spectrum of the posterior covariance being the key object to understanding the geometry throughout the diffusion process.

Algebra is the offer made by the devil to the mathematician. The devil says: ‘I will give you this powerful machine, it will answer any question you like. All you need to do is give me your soul: give up geometry and you will have this marvellous machine.’

— *Michael Atiyah*

1 Introduction

1.1 Initial Project

The progress in generative modelling over the last decade has been rapid. A key shift was the introduction of score based generative models (SGM), also known as diffusion models, which allow us to interpolate between a known distribution (usually a standard Gaussian), and a target distribution p_{data} , by training a network based on data $X \sim p_{\text{data}}$. By repeatedly adding noise to a target density, and training a denoising network, one can sample from p_{data} . Now, flow matching methods are the standard implementation of diffusion based generative models. These crucially can be viewed as a SGM.

While vast amounts of work has been done on improving the capabilities of diffusion models, the geometry of the entire diffusion process has been under explored.

Initially, this research internship aimed to investigate the geometry of both the autoencoder and denoiser in latent diffusion models. If one can show that the geometry of a target distribution influences the estimation of scores, discretisation error, or sampling stability, then the case can be made that when the autoencoder and diffusion model are trained jointly, that the autoencoder learns a more geometrically regular embedding.

We spent the internship so far almost entirely looking at the geometry of diffusion.

1.2 Existing work

There exists literature on how different geometric aspects affect the diffusion process, but these are mostly asymptotic results for highly regular target densities (Benton et al., 2024), (Potapchik et al., 2024), (Liang et al., 2025).

The main challenge of the question “how does geometry affect diffusion” is answering what we mean by “geometry”. There are some issues with existing methods which make use of dimension and reach/curvature of the support of p_{data} :

- Dimension:
 - Ill-posed to assign a dimension $d < D$ to a density which is absolutely continuous with respect to $\mathbb{R}^D$
 - Dimension of p_{data} may not be constant
 - Dimension of p_{data} may not be well defined everywhere, for example at intersections
 - Only considers dimension of target, now how effective dimension might change.
- Reach, curvature:
 - Can be $0, \infty$ at intersections or corners
 - Can vary wildly along the support of p_{data}

1.3 Objects Investigated

This gap in the literature is the primary of the project. As such, a new object other than dimension or reach or curvature had to be found. Throughout we investigated several before landing on the final and most promising object: the posterior covariance.

- Pullback metric of flow $G_{t,T}(x)$
- Level sets
- Ridges
- Posterior Covariance

[!!!make a table comparing them after meeting!!!]

1.4 Main Contribution

We find that the spectrum of the posterior covariance captures geometric variation, and use functionals of the spectrum to understand geometric variation throughout the diffusion process. One of the things this allows us to do is pick a discretisation schedule to reduce the discretisation error in sampling.

2 Background

2.1 The Manifold Hypothesis

The manifold hypothesis is an explanation of the effectiveness of machine learning for unstructured high-dimensional data. It states that a distribution $X \sim p_{\text{data}}, X \in \mathbb{R}^D$ lies on (or is concentrated around) a submanifold $\mathcal{M}$ of lower dimension $d \ll D$. This effectively makes the number of degrees of freedom d , allowing models to learn these lower dimensional spaces in settings where we may have $n < D$ data points.

The main argument for the manifold hypothesis is the curse of dimensionality, whereby there are too many dimensions making data very sparse and hard for models to learn from. Many counterintuitive behaviours are exhibited in high dimensional space, for example standard Gaussian data from $\mathcal{N}(0; \text{Id}_D)$ concentrating around the shell of radius $\sqrt{D}$.

The justification of the manifold hypothesis is that it is a good candidate model for naturally arising data, which tends to be locally continuous, and have fewer degrees of freedom. The space of natural images is far smaller than the space of all possible images, for example.

There are variations of the manifold hypothesis, with less restrictive assumptions. These include:

- The union of manifolds hypothesis ([Brown et al., 2022a](#)), ([Brown et al., 2022b](#)), which argue that standard manifold hypothesis is incomplete, and a union of manifolds of potentially different dimensions is a more realistic assumption
- The stratified space hypothesis ([Aamari and Berenfeld, 2024](#)), which extends this idea but allow for intersections between manifolds of different dimension

For completeness, we include here the definition of a manifold and the reach.

A d -dimensional manifold is a topological space $\mathcal{M}$ that is locally isomorphic to an open subset of $\mathbb{R}^d$. So, for each $y \in \mathcal{M}$, there is an open set $U_y \subseteq \mathcal{M}$ and a continuous function $\Phi_y : U_y \rightarrow \mathbb{R}^d$ such that Φ_y is a homeomorphism onto its image. The smoothness of the manifold is defined as the smoothness of the functions Φ_y .

When a manifold M is embedded into $\mathbb{R}^D$, to control the regularity of the embedding ([Federer, 1959](#)) introduce the notion of the reach $\tau = \tau(\mathcal{M})$ which is defined as

$$\tau := \sup \{ \varepsilon : \forall x \in \mathcal{M}^\varepsilon, \quad \exists! y \in \mathcal{M} \text{ s.t. } \text{dist}(x, \mathcal{M}) = \|x - y\| \},$$

where $\mathcal{M}^\varepsilon = \{x \in \mathbb{R}^D : \text{dist}(x, \mathcal{M}) < \varepsilon\}$ is the ε -neighborhood of $\mathcal{M}$. Equivalently, τ is the supremum over the radii of neighborhoods of $\mathcal{M}$ for which the projection is unique.

2.2 Flow Matching as a Score-Based Generative Model

Score-Based Generative Models (SGMs) (Ho et al., 2020; Song and Ermon, 2019; Song et al., 2020b) formulate generative modeling by estimating gradients of the log probability density (the score function). As noted by (Song and Ermon, 2019): “Because gradients can be ill-defined and hard to estimate when the data resides on **low-dimensional manifolds**, we perturb the data with different levels of Gaussian noise, and jointly estimate the corresponding scores, i.e., the vector fields of gradients of the perturbed data distribution for all noise levels.”

Flow Matching (Lipman et al., 2022; Albergo et al., 2025) provides a continuous-time vector field framework that unifies and generalises SGMs. Below, we outline the connection between Gaussian continuous flows and score matching.

Gaussian Flows and Interpolation Schedule. In generative modeling via flow matching, we aim to construct a path transforming a base noise distribution $p_1 = \mathcal{N}(0, I_d)$ into a target data distribution p_0 . Let $X_0 \sim p_0$ and $X_1 = \varepsilon \sim \mathcal{N}(0, I_d)$. The linear conditional flow $X_t = \varphi_t(X_0, \varepsilon)$ interpolates between data and noise according to:

$$X_t = c_t X_0 + \sigma_t \varepsilon \quad (1)$$

where $c_t, \sigma_t \in \mathbb{R}$ are differentiable schedule functions satisfying the boundary conditions $c_0 = 1, \sigma_0 = 0$ and $c_1 = 0, \sigma_1 = 1$. Differentiating with respect to time t yields the conditional velocity field:

$$\dot{X}_t = \dot{c}_t X_0 + \dot{\sigma}_t \varepsilon \quad (2)$$

Vector Field Formulation. computational marginal probability density p_t of X_t evolves according to the continuity equation $\partial_t p_t = -\nabla \cdot (v_t p_t)$. The marginal vector field $v_t(x)$ that generates this trajectory is given by the conditional expectation:

$$v_t(x) = \mathbb{E} [\dot{X}_t \mid X_t = x] = \dot{c}_t \mathbb{E} [X_0 \mid X_t = x] + \dot{\sigma}_t \mathbb{E} [\varepsilon \mid X_t = x] \quad (3)$$

Derivation via Tweedie’s Formula. Because p_t is the distribution of $c_t X_0$ corrupted by additive Gaussian noise $\sigma_t \varepsilon \sim \mathcal{N}(0, \sigma_t^2 I_d)$, Tweedie’s denoising identity provides closed-form relationships between conditional expectations and the score function $\nabla_x \log p_t(x)$:

$$\mathbb{E} [\varepsilon \mid X_t = x] = -\sigma_t \nabla_x \log p_t(x) \quad (4)$$

$$\mathbb{E} [X_0 \mid X_t = x] = \frac{x + \sigma_t^2 \nabla_x \log p_t(x)}{c_t} \quad (5)$$

Substituting equations (4) and (5) into the marginal velocity expression

yields:

$$\begin{aligned}
v_t(x) &= \dot{c}_t \left(\frac{x + \sigma_t^2 \nabla_x \log p_t(x)}{c_t} \right) + \dot{\sigma}_t (-\sigma_t \nabla_x \log p_t(x)) \\
&= \frac{\dot{c}_t}{c_t} x + \left(\frac{\dot{c}_t \sigma_t^2}{c_t} - \dot{\sigma}_t \sigma_t \right) \nabla_x \log p_t(x) \\
&= \frac{\dot{c}_t}{c_t} x + \sigma_t^2 \left(\frac{\dot{c}_t}{c_t} - \frac{\dot{\sigma}_t}{\sigma_t} \right) \nabla_x \log p_t(x)
\end{aligned} \tag{6}$$

Equivalence of Parameterisations. Equation (6) demonstrates that the velocity field $v_t(x)$ is completely determined by the schedules c_t, σ_t and the score function $\nabla_x \log p_t(x)$. Consequently, the following target parameterisations are equivalent up to linear transformations:

$$\nabla_x \log p_t(x) \iff \mathbb{E}[\varepsilon \mid X_t = x] \iff \mathbb{E}[X_0 \mid X_t = x] \iff v_t(x) \tag{7}$$

Sampling via the deterministic flow ODE $\dot{x}_t = v_t(x_t)$ parameterised by a score model recovers the probability flow ODE (e.g., DDIM (Song et al., 2020a)).

2.3 Tweedie’s 1st and 2nd Order Formulae

Consider the canonical Gaussian perturbation model $X_t = X_0 + \sigma_t \varepsilon$, where $\varepsilon \sim \mathcal{N}(0, I_d)$ and $X_0 \sim p_0$. The conditional density of X_t given $X_0 = x_0$ is:

$$p_t(x \mid x_0) = (2\pi\sigma_t^2)^{-d/2} \exp\left(-\frac{\|x - x_0\|^2}{2\sigma_t^2}\right) \tag{8}$$

The marginal density $p_t(x)$ is:

$$p_t(x) = \int p_t(x \mid x_0) p_0(x_0) dx_0 \tag{9}$$

2.3.1 First Derivative (Score Function)

Differentiating $p_t(x)$ under the integral sign:

$$\nabla_x p_t(x) = -\frac{1}{\sigma_t^2} \int (x - x_0) p_t(x \mid x_0) p_0(x_0) dx_0 \tag{10}$$

Dividing by $p_t(x)$ and applying Bayes’ rule $p(x_0 \mid x) = \frac{p_t(x \mid x_0) p_0(x_0)}{p_t(x)}$ yields Tweedie’s first-order formula:

$$\nabla_x \log p_t(x) = \frac{\nabla_x p_t(x)}{p_t(x)} = \frac{\mathbb{E}[X_0 \mid X_t = x] - x}{\sigma_t^2} \tag{11}$$

2.3.2 Second Derivative (Hessian Matrix/Score Jacobian)

Expanding the Hessian matrix of the log-marginal density:

$$\nabla_x^2 \log p_t(x) = \frac{\nabla_x^2 p_t(x)}{p_t(x)} - (\nabla_x \log p_t(x)) (\nabla_x \log p_t(x))^T \quad (12)$$

Differentiating $\nabla_x p_t(x)$ a second time:

$$\nabla_x^2 p_t(x) = \int \left(-\frac{1}{\sigma_t^2} I_d + \frac{(x - x_0)(x - x_0)^T}{\sigma_t^4} \right) p_t(x | x_0) p_0(x_0) dx_0 \quad (13)$$

Dividing by $p_t(x)$:

$$\frac{\nabla_x^2 p_t(x)}{p_t(x)} = -\frac{1}{\sigma_t^2} I_d + \frac{1}{\sigma_t^4} \mathbb{E}[(x - X_0)(x - X_0)^T | X_t = x] \quad (14)$$

2.3.3 Posterior Covariance Relation

Substituting (11) and (14) into (12):

$$\begin{aligned} \nabla_x^2 \log p_t(x) = & -\frac{1}{\sigma_t^2} I_d + \frac{1}{\sigma_t^4} \left(\mathbb{E}[(x - X_0)(x - X_0)^T | X_t = x] \right. \\ & \left. - \mathbb{E}[x - X_0 | X_t = x] \mathbb{E}[x - X_0 | X_t = x]^T \right) \end{aligned} \quad (15)$$

Since x is non-random given $X_t = x$, the bracketed term is the conditional covariance $\text{Cov}(X_0 | X_t = x)$. Thus, Tweedie's second-order formula simplifies to:

$$\nabla_x^2 \log p_t(x) = \frac{1}{\sigma_t^4} \text{Cov}(X_0 | X_t = x) - \frac{1}{\sigma_t^2} I_d \quad (16)$$

This identity will be used later in [section 6](#).

3 Anisotropic Stability Bound

We record a stability estimate for generator matching that keeps track of *which* directions of a drift error are amplified and which are dampened. The scalar one-sided Lipschitz estimate propagates every error with a single worst-case rate, thereby being sharp only in an isotropic sense. The estimate below replaces that scalar rate by the derivative flow of the learnt dynamics, a matrix that sees the direction of the error.

3.1 The deterministic transport mechanism

Consider two time-dependent velocity fields $v_t, \hat{v}_t : \mathbb{R}^D \rightarrow \mathbb{R}^D$ for $0 \leq t \leq T$, smooth enough in space and integrable in time that their characteristic flows are globally defined and differentiable in the initial condition:

$$\frac{d}{ds} \Phi_{t,s}(x) = v_s(\Phi_{t,s}(x)), \quad \frac{d}{ds} \hat{\Phi}_{t,s}(x) = \hat{v}_s(\hat{\Phi}_{t,s}(x)), \quad \Phi_{t,t} = \hat{\Phi}_{t,t} = \text{Id}. \quad (17)$$

We write $\delta v_t := v_t - \hat{v}_t$ for the velocity error and let v_t be the exact field and $\hat{v}_t$ the learnt one.

The estimate is built entirely from the derivative of the *learnt* flow, which is available to us from its network,

$$\hat{J}_{t,T}(x) := D_x \hat{\Phi}_{t,T}(x), \quad \frac{d}{ds} \hat{J}_{t,s}(x) = \nabla \hat{v}_s(\hat{\Phi}_{t,s}(x)) \hat{J}_{t,s}(x), \quad \hat{J}_{t,t}(x) = \text{Id}. \quad (18)$$

Thus $\hat{J}_{t,T}(x)h$ is the terminal displacement at time T produced by an infinitesimal perturbation h applied at (t, x) and transported by the learnt dynamics. An error at time t is not scaled by a worst-case constant but carried by $\hat{J}_{t,T}$ along its own direction. The following identity makes this exact.

Lemma 3.1 (Telescoping identity along characteristics). *For every initial point $x_0 \in \mathbb{R}^D$, with $X_t := \Phi_{0,t}(x_0)$ the exact characteristic,*

$$\Phi_{0,T}(x_0) - \hat{\Phi}_{0,T}(x_0) = \int_0^T \hat{J}_{t,T}(X_t) \delta v_t(X_t) dt. \quad (19)$$

Proof. Define $F(t) := \hat{\Phi}_{t,T}(X_t)$, which follows the exact dynamics up to time t and the learnt dynamics from t to T . Its endpoints are the fully learnt and fully exact terminal points, $F(0) = \hat{\Phi}_{0,T}(x_0)$ and $F(T) = \Phi_{0,T}(x_0)$. Differentiating in t ,

$$\frac{d}{dt} F(t) = \partial_t \hat{\Phi}_{t,T}(X_t) + D \hat{\Phi}_{t,T}(X_t) \dot{X}_t.$$

By starting the learnt flow an instant later we remove one short learnt step, so $\partial_t \hat{\Phi}_{t,T}(x) = -D \hat{\Phi}_{t,T}(x) \hat{v}_t(x)$. Together with $\dot{X}_t = v_t(X_t)$ and $D \hat{\Phi}_{t,T}(X_t) = \hat{J}_{t,T}(X_t)$ this gives

$$\frac{d}{dt} F(t) = \hat{J}_{t,T}(X_t) (v_t(X_t) - \hat{v}_t(X_t)) = \hat{J}_{t,T}(X_t) \delta v_t(X_t).$$

Integrating $F(T) - F(0)$ from 0 to T yields (19). $\square$

Identity (19) highlights that the error made at each time enters the terminal gap through the derivative flow evaluated in the actual direction of that error, with no worst-case constant and no loss.

3.2 From the telescoping identity to the matrix bound

We turn the pathwise identity into a bound on the terminal laws by coupling. Let $X_0 \sim \rho_0$ and transport it by both flows, $X_t := \Phi_{0,t}(X_0)$ and $\hat{X}_t := \hat{\Phi}_{0,t}(X_0)$, with terminal laws $\rho_T := \text{Law}(X_T)$ and $\hat{\rho}_T := \text{Law}(\hat{X}_T)$. Running both flows from the same X_0 is a coupling of ρ_T and $\hat{\rho}_T$, so $W_1(\rho_T, \hat{\rho}_T) \leq \mathbb{E}\|X_T - \hat{X}_T\|$. Applying Lemma 3.1 pathwise at $x_0 = X_0$ and the triangle inequality,

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \mathbb{E} \|\hat{J}_{t,T}(X_t) \delta v_t(X_t)\| dt = \int_0^T \int_{\mathbb{R}^D} \|\hat{J}_{t,T}(x) \delta v_t(x)\| \rho_t(dx) dt, \quad (20)$$

where the last step uses $X_t \sim \rho_t$. It remains only to name the quadratic form hidden in the norm.

Theorem 3.2 (Anisotropic W_1 stability). *Define the pullback amplification metric*

$$G_{t,T}(x) := \hat{J}_{t,T}(x)^\top \hat{J}_{t,T}(x), \quad \text{so that} \quad \|\hat{J}_{t,T}(x)h\|^2 = \langle h, G_{t,T}(x)h \rangle. \quad (21)$$

Then

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \int_{\mathbb{R}^D} \langle \delta v_t(x), G_{t,T}(x) \delta v_t(x) \rangle^{1/2} \rho_t(dx) dt. \quad (22)$$

Proof. Substitute the identity in (21) into (20). $\square$

The matrix $G_{t,T}(x)$ is the Euclidean metric at the terminal time pulled back to the tangent space at (t, x) by the learnt flow; it lives at time t but measures squared lengths after transport to T . Its eigenvectors are the input directions at (t, x) and its eigenvalues are the squared stretching factors of those directions. In the deterministic case (21) is an exact identity, so no Cauchy–Schwarz slack is incurred in passing from the derivative flow to the quadratic form. The bound (22) is anisotropic precisely because it evaluates $G_{t,T}$ on the actual error direction $\delta v_t(x)$ rather than on the worst possible direction.

3.3 Recovering the one-sided Lipschitz estimate

The classical scalar estimate is the isotropic specialisation of Theorem 3.2, obtained by discarding the direction of the error. Let $S_t(x) := \frac{1}{2}(\nabla \hat{v}_t(x) + \nabla \hat{v}_t(x)^\top)$ be the symmetric part of the learnt Jacobian, and suppose the learnt field is one-sided Lipschitz with rate ℓ_t , meaning

$$\lambda_{\max}(S_t(x)) \leq \ell_t \quad \text{for all } x. \quad (23)$$

The condition constrains expansion by only bounding the largest rate at which two trajectories separate, and places no bound on contraction or on the rotational part of $\nabla \hat{v}_t$.

Corollary 3.3 (Scalar one-sided Lipschitz bound). *Under (23),*

$$G_{t,T}(x) \preceq \exp\left(2 \int_t^T \ell_s \, ds\right) \text{Id}, \quad (24)$$

and consequently

$$W_1(\rho_T, \hat{\rho}_T) \leq \int_0^T \exp\left(\int_t^T \ell_s \, ds\right) \|\delta v_t\|_{L^1(\rho_t)} \, dt. \quad (25)$$

Proof. Fix a direction h and set $z_s := \hat{J}_{t,s}(x)h$, which solves $\dot{z}_s = \nabla \hat{v}_s(\hat{\Phi}_{t,s}(x)) z_s$ by (18). Only the symmetric part contributes to the growth of the norm, so

$$\frac{d}{ds} \|z_s\|^2 = 2 \langle z_s, S_s(\hat{\Phi}_{t,s}(x)) z_s \rangle \leq 2 \ell_s \|z_s\|^2.$$

Grönwall's lemma gives $\|\hat{J}_{t,T}(x)h\| \leq \exp\left(\int_t^T \ell_s \, ds\right) \|h\|$, which is (24). Inserting (24) into (22) and using $\langle \delta v_t, G_{t,T} \delta v_t \rangle^{1/2} \leq \exp\left(\int_t^T \ell_s \, ds\right) \|\delta v_t\|$ yields (25). $\square$

Corollary 3.3 exhibits the scalar estimate as the Loewner domination (24): the matrix $G_{t,T}$ is replaced by its largest eigenvalue, uniformly over directions. While bound is correct, it does not make use of the eigenstructure of $G_{t,T}$, and in particular treats a strongly contracted error direction as though it were the most expansive one. The gap between the scalar and anisotropic bounds can be exacerbated when errors lie mostly in the normal direction to a target manifold, as we will see.

3.4 A two-direction picture

The gap between the matrix and scalar bounds is already visible in the simplest anisotropic model. Fix an orthogonal splitting $\mathbb{R}^D = E_{\parallel} \oplus E_{\perp}$ into a tangent-like block and a normal-like block, write $x = (x_{\parallel}, x_{\perp})$ accordingly, and take the learnt velocity to be linear, block-diagonal, expansive on $E_{\parallel}$ and contractive on $E_{\perp}$:

$$\hat{v}_t(x_{\parallel}, x_{\perp}) = (\ell_t x_{\parallel}, -\gamma_t x_{\perp}), \quad \gamma_t \geq 0. \quad (26)$$

Its Jacobian $\nabla \hat{v}_t = \text{diag}(\ell_t I_{\parallel}, -\gamma_t I_{\perp})$ is constant in space, so the derivative flow (18) integrates blockwise and the pullback metric (21) is diagonal:

$$\hat{J}_{t,T} = \begin{pmatrix} e^{\int_t^T \ell_s \, ds} I_{\parallel} & 0 \\ 0 & e^{-\int_t^T \gamma_s \, ds} I_{\perp} \end{pmatrix}, \quad G_{t,T} = \begin{pmatrix} e^{2 \int_t^T \ell_s \, ds} I_{\parallel} & 0 \\ 0 & e^{-2 \int_t^T \gamma_s \, ds} I_{\perp} \end{pmatrix}. \quad (27)$$

Splitting the error as $\delta v_t = \delta v_t^{\parallel} + \delta v_t^{\perp}$ with $\delta v_t^{\parallel} \in E_{\parallel}$ and $\delta v_t^{\perp} \in E_{\perp}$, the integrand of the matrix bound (22) separates,

$$\langle \delta v_t, G_{t,T} \delta v_t \rangle^{1/2} = \left(e^{2 \int_t^T \ell_s \, ds} \|\delta v_t^{\parallel}\|^2 + e^{-2 \int_t^T \gamma_s \, ds} \|\delta v_t^{\perp}\|^2 \right)^{1/2}. \quad (28)$$

A tangent error is carried to the terminal time by the expansion factor $e^{\int_t^T \ell}$; a normal error is carried by the contraction factor $e^{-\int_t^T \gamma}$ and is therefore damped.

The scalar estimate cannot distinguish the two. Its rate is the largest eigenvalue of the symmetric Jacobian,

$$\lambda_{\max}(S_t) = \max\{\ell_t, -\gamma_t\} = \ell_t \quad (\ell_t \geq 0), \quad (29)$$

so Corollary 3.3 propagates *both* error components with the tangent expansion factor $e^{\int_t^T \ell}$. It replaces the normal contraction $e^{-\int_t^T \gamma}$ by expansion, overestimating the contribution of a normal error by the factor $e^{\int_t^T (\ell_s + \gamma_s) ds}$. When the normal block is strongly contracting (which is the case near a terminal manifold, where $\gamma_t \sim (T-t)^{-1}$ and $e^{-\int_t^T \gamma}$ decays polynomially in $T-t$) this overestimate is the difference between a persistent error and a damped one.

3.5 What the estimate requires of the error

The two bounds depend on the error through different amounts of information, and this raises a natural objection. The scalar bound (25) uses only the magnitude $\|\delta v_t\|_{L^1(\rho_t)}$. The matrix bound (22) uses the quadratic form $\langle \delta v_t, G_{t,T} \delta v_t \rangle$, that is, the direction of δv_t relative to the eigenframe of $G_{t,T}$. Since the exact field v_t is unknown, the error $\delta v_t = v_t - \hat{v}_t$ is never observed, and its direction even less so. Therefore, we conclude that this matrix bound, while tighter, is not any more practical than the scalar bound.

Despite observing that δv_t was usually aligned with v_t , and that in different instances of training δv_t seemed to be Gaussian, not distributional behaviour of δv_t was understood.

One assumption is that the projectors defining the frame (P_t onto the tangent block and $Q_t = I - P_t$ onto the normal block) must be a single field, determined by the target and shared across all learnt fields. Re-deriving the frame from each model's own error would reintroduce model dependence into the weights and make comparisons between models meaningless. We therefore fix P_t once, from the target density geometry, and read every error in that frame. So P_t is built from the spectrum of the posterior-covariance operator A_t in section 6.

4 Level Sets

As we smooth data which lies on (or is concentrated around) a manifold two main affects occur: smoothing in the normal direction, and smoothing and tangent direction. Smoothing in the tangent direction makes a density more uniform along a manifold, therefore our level sets should be better behaved. Here we investigate how level sets evolve.

Let $p^\star : \mathbb{R}^D \rightarrow \mathbb{R}_+$ denote the initial density, and let p_σ denote its Gaussian-smoothed version,

$$p_\sigma(x) = (G_\sigma * p^\star)(x), \quad G_\sigma(x) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right). \quad (30)$$

For a regular value c , the corresponding level set is

$$\mathcal{M}_c^\sigma := \{x \in \mathbb{R}^D : p_\sigma(x) = c\}, \quad \nabla p_\sigma(x) \neq 0 \quad \text{for } x \in \mathcal{M}_c^\sigma. \quad (31)$$

The unit normal to $\mathcal{M}_c^\sigma$ is

$$n(x) = \frac{\nabla p_\sigma(x)}{\|\nabla p_\sigma(x)\|}. \quad (32)$$

For a tangent direction $v \in T_x \mathcal{M}_c^\sigma$, the derivative of the normal is

$$D_v n = \frac{1}{\|\nabla p_\sigma\|} (I - nn^\top) D^2 p_\sigma v = \frac{1}{\|\nabla p_\sigma\|} P_T D^2 p_\sigma v, \quad (33)$$

where

$$P_T := I - nn^\top \quad (34)$$

is the orthogonal projection onto the tangent space $T_x \mathcal{M}_c^\sigma$. Thus, the variation of the normal field, and hence the local curvature of the level set, is determined by the Hessian of the density.

Gaussian smoothing regularises not only the density itself, but all of its spatial derivatives. In particular, for any $k \geq 0$,

$$D^k p_\sigma = D^k (G_\sigma * p^\star) = (D^k G_\sigma) * p^\star. \quad (35)$$

Hence increasing σ smooths the density and simultaneously suppresses fine-scale variation in its gradient, Hessian, and higher-order derivatives.

This smoothing has a particularly simple interpretation in the Fourier domain. Writing $\hat{p}_\sigma$ for the Fourier transform,

$$\hat{p}_\sigma(\xi) = \hat{G}_\sigma(\xi) \hat{p}^\star(\xi) = \exp\left(-\frac{\sigma^2 \|\xi\|^2}{2}\right) \hat{p}^\star(\xi). \quad (36)$$

Thus a Fourier mode of spatial frequency $\|\xi\|$ is attenuated by $\exp(-\sigma^2 \|\xi\|^2/2)$. High-frequency components decay quadratically faster in frequency. Consequently, Gaussian smoothing preferentially removes the fine-scale spatial variation of $p^\star$.

The effect on the geometry of a level set follows directly from (33). Its shape operator is

$$S := -P_T Dn P_T = -\frac{1}{\|\nabla p_\sigma\|} P_T D^2 p_\sigma P_T, \quad (37)$$

up to the choice of orientation for the normal. The eigenvalues of S are the principal curvatures of $\mathcal{M}_c^\sigma$. Therefore, Gaussian smoothing affects the shape of the level set through two quantities: the Hessian $D^2 p_\sigma$, which controls variation of the normal, and $\|\nabla p_\sigma\|$, which sets the scale relating density variation to geometric variation. In particular, since Gaussian smoothing suppresses high-frequency components of $D^2 p_\sigma$, it suppresses the fine-scale variation of the shape operator and hence removes small-scale geometric structure from the level set.

It is important to distinguish this suppression of fine-scale variation from a pointwise reduction in curvature. In general, Gaussian smoothing does not imply that the principal curvatures of a fixed level set decrease monotonically with σ . Indeed, the shape operator

$$S_\sigma(x) = -\frac{1}{\|\nabla p_\sigma(x)\|} P_T D^2 p_\sigma(x) P_T, \quad x \in \mathcal{M}_c^\sigma, \quad (38)$$

depends on both $D^2 p_\sigma$ and ∇p_σ , and is evaluated on the moving level set $\mathcal{M}_c^\sigma$. Although Gaussian smoothing suppresses high-frequency components of $D^2 p_\sigma$, it simultaneously changes ∇p_σ and the location of the level set itself. Hence there is no general pointwise monotonicity of S_σ or of its eigenvalues.

Moreover, the topology of $\mathcal{M}_c^\sigma$ can change with σ . Such a change occurs when the level c passes through a critical point, i.e., when

$$p_\sigma(x) = c, \quad \nabla p_\sigma(x) = 0. \quad (39)$$

For example, two components of a superlevel set may merge through a neck as σ increases. Thus, while smoothing removes fine-scale spatial variation from the density and its derivatives, it should not be interpreted as a monotone decrease of the curvature of every individual level set. Instead that Gaussian smoothing preferentially suppresses spatial frequencies at scales smaller than σ , and therefore suppresses fine-scale geometric variation encoded by the derivatives of the density.

More generally, adding Gaussian noise monotonically suppresses all L^2 Sobolev energies of the smoothed density. Let

$$p_\sigma = G_\sigma * p^*, \quad G_\sigma(x) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left(-\frac{\|x\|^2}{2\sigma^2}\right), \quad (40)$$

where D is the ambient dimension. Since $\partial_{\sigma^2} p_\sigma = \frac{1}{2} \Delta p_\sigma$, we have, for every integer $k \geq 0$,

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int_{\mathbb{R}^D} |D^k p_\sigma(x)|^2 dx = -\frac{1}{2} \int_{\mathbb{R}^D} |\nabla D^k p_\sigma(x)|^2 dx \leq 0. \quad (41)$$

Thus, increasing the noise level σ monotonically decreases the L^2 norms of all spatial derivatives of p_σ . In particular,

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int |\nabla p_\sigma|^2 dx = -\frac{1}{2} \int |\Delta p_\sigma|^2 dx \leq 0, \quad (42)$$

$$\frac{\partial}{\partial \sigma^2} \frac{1}{2} \int |D^2 p_\sigma|^2 dx = -\frac{1}{2} \int |\nabla \Delta p_\sigma|^2 dx \leq 0. \quad (43)$$

Equivalently, in the Fourier domain,

$$\frac{1}{2} \int_{\mathbb{R}^D} |D^k p_\sigma(x)|^2 dx \propto \frac{1}{2} \int_{\mathbb{R}^D} \|\xi\|^{2k} e^{-\sigma^2 \|\xi\|^2} |\widehat{p^\star}(\xi)|^2 d\xi. \quad (44)$$

Hence increasing σ progressively removes increasingly fine-scale structure, with spatial frequency $\|\xi\|$ suppressed according to $e^{-\sigma^2 \|\xi\|^2}$.


This relates to the conventional wisdom where more refined time steps are required as we approach closer to our final density to capture the higher frequency structure of density. This is particularly apparent for image and audio data. JPEG compression takes advantage of writing the image as a Fourier decomposition, and in doing so discards the high frequency terms which are not noticeable to the human eye.

5 Ridges


5.1 Motivation

An at first unintuitive aspect of the manifold hypothesis is that the lower dimensional space around which data concentrate is not necessarily, or in fact usually, the initial manifold. This is because as the amount of noise that we add increases the effective radius with which we convolve grows and includes more of our initial manifold. We see a noise density concentrating around a different space. This begs the question “what is this new space and how does it behave and evolve?”

This becomes immediately obvious when considering smoothing a pair of distinct Bernoulli random variables: the modes of the sum is not the union of the modes.



(a) Gaussian smoothing moves modes in 1-D



(b) Gaussian smoothing moves the ridge in 2-D

Figure 1: Comparison of ridge structures.

With this intuition it becomes fairly easy to rediscover the definition of a ridge. It is where the normal component of our changing density is zero. But a priori we do not know which components are normal.

The resolution is to let the density define its own normal directions. At a point on a d -dimensional ridge the density behaves like a mountain ridge: it is at a local maximum in the $D - d$ directions transverse to the ridge, while along the remaining d directions it varies slowly. Transverse to the ridge the density curves sharply downward; along it, the curvature is mild. This is a statement about the Hessian. Let $g(x) = \nabla p(x)$ and $H(x) = \nabla^2 p(x)$, and order the eigenvalues of $H(x)$ as

$$\lambda_1(x) \geq \dots \geq \lambda_d(x) \geq \lambda_{d+1}(x) \geq \dots \geq \lambda_D(x),$$

with orthonormal eigenvectors $u_1(x), \dots, u_D(x)$. The directions of sharpest

downward curvature are those with the most negative eigenvalues, so we define the normal space at x to be the span of the bottom $D - d$ eigenvectors, and the tangent space to be the span of the top d . Writing $V(x) = [u_{d+1}(x) \mid \cdots \mid u_D(x)]$ for the normal frame, $L(x) = V(x)V(x)^\top$ is the projector onto the normal space. This is the point of the construction: which directions are normal is not supplied from outside, it is read off from the second-order structure of the smoothed density itself.

A ridge point is then one from which the density cannot be increased by moving normally. Equivalently, the normal component of the gradient vanishes. Define the projected gradient

$$G(x) = L(x)g(x),$$

and the ridge as the set where it is zero and the transverse curvature is genuinely downward:

$$R = \{x : \|G(x)\| = 0, \lambda_{d+1}(x) < 0\}.$$

The second condition, $\lambda_{d+1}(x) < 0$, requires the largest of the normal eigenvalues to still be negative, so that every normal direction is a direction of local descent and the point is a maximum transverse to the ridge rather than a saddle. The case $d = 0$ recovers ordinary modes: $L(x) = I$, the condition $\|G(x)\| = 0$ becomes $\nabla p(x) = 0$, and $\lambda_1(x) < 0$ asks for a local maximum. A d -ridge relaxes this by demanding a critical point only in the $D - d$ normal directions, and negative curvature only in the corresponding block.

5.2 The score concentrates on the ridge

We first record a convenient fact: the ridge of $\log p_\sigma$ and the ridge of p_σ are the same set. Since $\log$ is strictly increasing, $\nabla \log p_\sigma = p_\sigma^{-1} \nabla p_\sigma$ vanishes exactly where ∇p_σ does, and the two Hessians differ by the rank-one term $p_\sigma^{-1} \nabla p_\sigma \nabla p_\sigma^\top$, which vanishes wherever the gradient does. On the ridge gradient, projected gradient and normal eigenframe all agree between p_σ and $\log p_\sigma$. We work with $\log p_\sigma$ from now on, for two reasons: it is the object the generative model actually learns, its gradient being the score; and as shown in (Genovese et al., 2014), the surrogate property Theorem 5.1 holds in an $O(1)$ neighbourhood of $\mathcal{M}$ for the log-density but only in an $O(\sigma)$ neighbourhood for the density itself.

Theorem 5.1 (Surrogate Ridge Approximation, Theorem 7 of (Genovese et al., 2014)). *Consider the Hidden Manifold Model with small noise variance σ^2 , and assume that the manifold $\mathcal{M}$ has no boundary. Then the hyper-ridge set R of the density p satisfies*

$$\text{Haus}(\mathcal{M}, R) = O\left(\sigma^2 \log \frac{1}{\sigma}\right),$$

and R is topologically homeomorphic to $\mathcal{M}$. Consequently, as the noise parameter $\sigma \rightarrow 0$ with $\mathcal{M}$ held fixed, the ridge R converges to $\mathcal{M}$.

Inside the tube of radius $O(\sigma)$ around the manifold, the score is overwhelmingly a normal restoring field. (Genovese et al., 2014) shows that $\nabla \log p_\sigma(x) = -\frac{1}{\sigma^2}((x - \hat{x}) + O(K_\sigma^2))$, where $\hat{x} = \pi_{\mathcal{M}}(x)$ is the projection onto $\mathcal{M}$ and

$$K_\sigma = \sqrt{2\sigma^2 \log(\sigma^{-(A+D)})} = O\left(\sigma \sqrt{\log \frac{1}{\sigma}}\right).$$

The normal component has size $O(\sigma^{-1})$ while the tangential component is $O(1)$: as $\sigma \rightarrow 0$ the score is, to leading order, a restoring force. And this restoring force is towards the ridge, not $\mathcal{M}$

Writing the score as a first-order expansion about the manifold projection $m = \pi_{\mathcal{M}}(x)$ and, about the ridge projection $r = \pi_{R_\sigma}(x)$:

$$\nabla \log p_\sigma(x) = \nabla_{\mathcal{M}} \log p_\sigma(m) - \frac{1}{\sigma^2}(x - m) + O(K_\sigma^2) + o\left(\frac{\|x - m\|}{\sigma^2}\right), \quad (45)$$

$$\nabla \log p_\sigma(x) = \nabla_{R_\sigma} \log p_\sigma(r) - \frac{1}{\sigma^2}(x - r) + o\left(\frac{\|x - r\|}{\sigma^2}\right). \quad (46)$$

Expanding about the manifold leaves a residual $O(K_\sigma^2)$, while expanding about the ridge does not. This residual is a constant normal bias of magnitude $K_\sigma^2/\sigma^2 = O(\log \frac{1}{\sigma})$, present even on the manifold $\mathcal{M}$ itself. The zero of the normal restoring field has been pushed off $\mathcal{M}$ by exactly $O(K_\sigma^2)$ in the normal direction. That displaced zero is the ridge, and $O(K_\sigma^2)$ is precisely the Hausdorff bias $\text{Haus}(\mathcal{M}, R_\sigma) = O(\sigma^2 \log \frac{1}{\sigma})$ of Theorem 5.1.

Projecting (46) onto the normal space N_r at the ridge point makes this exact. The intrinsic gradient $\nabla_{R_\sigma} \log p_\sigma(r)$ lies in $T_r R_\sigma$ and is annihilated by $P_N(r)$, leaving

$$P_N(r) \nabla \log p_\sigma(x) = -\frac{1}{\sigma^2}(x - r) + o\left(\frac{\|x - r\|}{\sigma^2}\right).$$

The normal score is a pure linear restoring field toward the ridge, with slope $-\sigma^{-2}$ and no curvature-induced remainder.

This also reinforces our definition of the ridge as the locus where the normal component of the smoothed gradient vanishes. We see that this locus is the genuine attractor of the score's dominant term.

5.3 Failings of the Ridge

The objects the ridge is built from are smooth. For $t > 0$ the density p_t is a Gaussian convolution, hence real-analytic. Thus the posterior mean m_t , the covariance field A_t , and the Hessian $\nabla^2 \log p_t = (A_t - I)/t$ are all analytic in y , and their eigenvalues $\lambda_1(y) \geq \dots \geq \lambda_D(y)$ vary continuously everywhere (see section 6). Yet the ridge set R_t is not a continuous functional of the data.

Two independent failure modes break the continuity of R_t :

1. **Subspace Ambiguity:** Where $\lambda_d - \lambda_{d+1} \rightarrow 0$, the normal projector P_N is ill-defined, rendering the condition $\|P_N \nabla \log p_t\| = 0$ meaningless as the normal frame loses its orientation.

2. Topological Jumps: Where $\lambda_{d+1}(y) = 0$, the concavity condition triggers, causing R_t to abruptly gain or lose entire connected components.

(Genovese et al., 2014) accounts for this by establishing stability under an assumed lower bound β on this eigengap. The instability belongs to the mapping $(m_t, A_t) \mapsto R_t$ itself, and this is expected since passing smooth analytic fields through rigid scalar thresholds inevitably creates discontinuities.

A second, more basic problem is that the split requires the integer d to be fixed in advance, and there is no consistent choice. The dimension is an input to the eigenspace selection, yet the natural scale-dependent notion of dimension, $d_{\text{eff}}(t) = \mathbb{E} \text{tr } A_t$, is neither an integer nor constant in t : it interpolates between the intrinsic dimension at small noise and zero at large noise, and it overshoots in between (Chapter 6).

Example of Discontinuity of the Ridge. Consider a distribution on $\mathbb{S}^1 \subset \mathbb{R}^2$ which is a 1:3 mixture of 2 uniform densities on 2 semicircles. We see that a small change in noise can drastically change the location and topology of the ridge. Additionally, knowing the ridge does not tell us how the data is distributed along it, and density might be negligible at points on the ridge.




Figure 2: Comparison of circle ridges at different noise levels σ , with initial density 1:3 mixture of 2 uniform densities on left and right semicircles, respectively.

While the ridge R_σ has desirable qualities for small σ , and is the natural attractor of the score, for larger σ it may be ill defined and can be discontinuous. Recalling that the ridge is a functional of the spectrum of the Jacobian of the score leads us to section 6.

6 Posterior Covariance Field

Here we introduce the ultimate object we analyse: posterior covariance fields.

Definition 6.1. Let $Y_t = X + \sqrt{t}Z$, $m_t(y) := \mathbb{E}[X \mid Y_t = y]$, and

$$A_t(y) := \frac{1}{t} \text{Cov}(X \mid Y_t = y).$$

We call the weighted object $\mathfrak{F}_t := \text{Law}(Y_t, m_t(Y_t), A_t(Y_t))$ the *posterior covariance field*.

There are a few reasons why it makes sense to look at this object. First, it is not plagued by the issues we encountered in the analysis of ridges or level sets in [section 5](#) and [section 4](#). These objects suffered from being too geometric for our use case, where we care about noise scales much greater than 0. We need a more analytic quantity, and the posterior covariance field is defined for all noise scales $t > 0$. In fact level sets and ridges can just be viewed as functionals of $\mathfrak{F}_t$.

Second, the posterior covariance captures the local geometry of the score. The posterior covariance corresponds to the Jacobian of the score. This link is given in [section 2](#). This not at all obvious fact is where we reap the rewards. Here we normalise the posterior covariance by t to account for the scale of noise added.

Third, here we deal with expectations, rather than L^∞ -type quantities like the reach τ_M . If there exists a region of the density which is geometrically complex and as a consequence has poor regularity, but which carries almost no mass, we should not let it dramatically change our understanding of the behaviour of the score.

6.1 Three functionals of the spectrum

All three quantities below are moments of the same object: the spectral measure of A_t , averaged over $y \sim p_t$. Write $a_1(y), \dots, a_D(y)$ for the eigenvalues of $A_t(y)$ and let $\mu_t := \text{Law}(a_I(Y_t))$ with I uniform on $\{1, \dots, D\}$, so that $\mathbb{E} \sum_i f(a_i) = D \int f d\mu_t$ for any f .

6.1.1 Effective dimension

Definition 6.2.

$$d_{\text{eff}}(t) := \mathbb{E}[\text{tr } A_t(Y_t)] = \mathbb{E} \sum_i a_i. \quad (47)$$

Proposition 6.3 (Exact MMSE identity). *Let $\text{MMSE}(t) := \mathbb{E}|X - m_t(Y_t)|^2$. Then*

$$d_{\text{eff}}(t) = \frac{\text{MMSE}(t)}{t}. \quad (48)$$

Proof. $\text{tr Cov}(X \mid Y_t) = \mathbb{E}[|X - m_t(Y_t)|^2 \mid Y_t]$; take expectations and divide by t . $\square$

Note in identity (48) $\text{MMSE}(t)$ is precisely the population value of the ε -prediction (or x_0 -prediction) denoising loss at noise level t , so it requires no computation beyond the training objective already being evaluated.

Link to dimension. An orthogonal projector of rank k has spectrum $\{1^{(k)}, 0^{(D-k)}\}$ and trace k . In the ideal manifold limit $A_t \approx P_{T_x \mathcal{M}}$ and hence $d_{\text{eff}}(t) \approx d$. Equivalently, via (48): a d -dimensional manifold leaves $O(t)$ posterior uncertainty in each of d tangent directions and $o(t)$ in the normal directions, so $\text{MMSE} \approx d \times t$. Each eigenvalue may be read as a fractional dimension: $a_i = 0.4$ contributes 0.4 effective directions.

Failure mode. d_{eff} is a first moment and therefore cannot distinguish $\text{diag}(1, 1, 1, 0, 0, 0)$ from $\text{diag}(\frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2}, \frac{1}{2})$: both have trace 3. It also overcounts when $a_i > 1$, in which case the “fractional dimension” reading breaks down entirely. The next two functionals repair these two deficiencies.

6.1.2 Distance to projector

Definition 6.4.

$$\Delta_{\text{proj}}(t) := \mathbb{E} \sum_i \min\{a_i^2, (a_i - 1)^2\}. \quad (49)$$

The summand is the squared distance from a_i to the set $\{0, 1\}$, and the two branches exchange at $a_i = \frac{1}{2}$:

$$\min\{a^2, (a - 1)^2\} = \begin{cases} a^2, & a \leq \frac{1}{2} \quad (\text{read as normal-like}), \\ (a - 1)^2, & a \geq \frac{1}{2} \quad (\text{read as tangent-like}). \end{cases} \quad (50)$$

Proposition 6.5. Let $\mathcal{P} = \{P \in \mathbb{R}^{D \times D} : P = P^\top = P^2\}$ be the set of orthogonal projectors and let A be symmetric with eigenvalues a_i . Then

$$\min_{P \in \mathcal{P}} \|A - P\|_F^2 = \sum_i \min\{a_i^2, (a_i - 1)^2\}, \quad \text{hence} \quad \Delta_{\text{proj}}(t) = \mathbb{E}[\text{dist}_F^2(A_t(Y_t), \mathcal{P})]. \quad (51)$$

Proof. $\mathcal{P}$ is invariant under conjugation by orthogonal matrices, and for symmetric A the Frobenius distance to a conjugation-invariant set is minimised by a P sharing the eigenvectors of A (a consequence of the Hoffman–Wielandt inequality). It then remains to round each a_i to the nearer of 0 or 1. $\square$

Thus Δ_{proj} measures how far, on average, is the soft posterior geometry from being a genuine tangent-space projector. It vanishes exactly when $A_t(Y_t)$ is a projector almost surely.

6.1.3 Ambiguity

Definition 6.6. With $(x)_+ := \max\{x, 0\}$,

$$\mathcal{A}(t) := \mathbb{E} \sum_i (a_i - 1)_+^2. \quad (52)$$

The motivation is dynamical rather than geometric. Let (a_i, v_i) be an eigenpair of A_t . $a_i > 1$ occurs if and only if the reverse probability-flow ODE is locally expansive along v_i . The threshold where posterior variance exceeds the injected noise variance:

$$a_i = \frac{\text{posterior variance along } v_i}{\text{noise variance } t} \quad \begin{cases} < 1 : & \text{denoising, contraction,} \\ = 1 : & \text{neutral,} \\ > 1 : & \text{amplification, expansion.} \end{cases} \quad (53)$$

The three design choices are then: threshold at 1 (only expansion is penalised), measure the *excess* $a_i - 1$ (the neutral point is 1, not 0), and square it (making the functional comparable to a squared Frobenius norm and penalising strong expansion superlinearly).

Link to ambiguity. Suppose the posterior at y is approximately $\frac{1}{2}\mathcal{N}(\mu_1, \Sigma) + \frac{1}{2}\mathcal{N}(\mu_2, \Sigma)$, corresponding to two competing clean configurations. Then

$$\text{Cov}(X \mid Y_t = y) = \Sigma + \frac{1}{4}(\mu_1 - \mu_2)(\mu_1 - \mu_2)^\top, \quad (54)$$

and the between-mode term can exceed tI along $\mu_1 - \mu_2$, giving $a_i > 1$. So $a_i > 1$ signals that the noisy observation fails to identify a single locally stable latent configuration.

Remark 6.7 (Scope of interpretation). Strictly, $a_i > 1$ is a diagnostic that posterior variance exceeds the noise scale. Multimodality is the typical, but not the only, mechanism: heavy-tailed unimodal p_0 (e.g. Laplace, Student- t) also produces $a_i > 1$. “Ambiguity” is the geometric reading of the phenomenon, not a theorem.

Remark 6.8 (Mean versus tail). $\mathcal{A}$ is an average over $y \sim p_t$ and can be small while $\sup_y a_{\max}(y)$ is large, if the expansive region carries little mass. In practice one should report both $\mathcal{A}(t)$ and a high quantile of $(a_{\max} - 1)_+$, since sampler failure is driven by trajectories that enter the expansive region, not by its measure.

6.2 Evolution law of MMSE

By [Theorem 6.9](#), differentiating $\text{MMSE}(t)$ under the Gaussian channel model $Y_t = X + \sqrt{t}Z$ gives:

$$\text{MMSE}'(t) = \frac{1}{t^2} \mathbb{E} \left[\text{tr} \left(\text{Cov}(X \mid Y_t)^2 \right) \right] \quad (55)$$

By definition, $A_t(y) = \frac{1}{t} \text{Cov}(X \mid Y_t = y)$. Squaring this relation and taking the trace yields:

$$\text{tr}(A_t(Y_t)^2) = \frac{1}{t^2} \text{tr}(\text{Cov}(X \mid Y_t)^2) = \|A_t(Y_t)\|_F^2 \quad (56)$$

Substituting this directly into the expression for $\text{MMSE}'(t)$:

$$\text{MMSE}'(t) = \mathbb{E}[\text{tr}(A_t(Y_t)^2)] = \mathbb{E}\|A_t(Y_t)\|_F^2 \quad (57)$$

6.3 MMSE-driven discretisation

We consider the additive Gaussian perturbation

$$X_t = X_0 + \sqrt{t} Z, \quad Z \sim \mathcal{N}(0, I_D), \quad (58)$$

so that, in the notation of the diffusion model,

$$c_t = 1, \quad \sigma_t^2 = t. \quad (59)$$

Let

$$m_t(x) := \mathbb{E}[X_0 \mid X_t = x] \quad (60)$$

denote the posterior mean and define the minimum mean-squared error

$$\text{MMSE}(t) := \mathbb{E}[\|X_0 - m_t(X_t)\|^2]. \quad (61)$$

Following the effective-dimension viewpoint, we write

$$d_{\text{eff}}(t) := \frac{\text{MMSE}(t)}{t}, \quad (62)$$

and hence

$$\text{MMSE}(t) = t d_{\text{eff}}(t). \quad (63)$$

Discretisation error. For additive Gaussian noise, Tweedie's formula gives

$$s_t(x) = \frac{m_t(x) - x}{t}. \quad (64)$$

The same martingale argument used in (Potapchik et al., 2024) to control the score variation in the discretised reverse process then yields, for two noise levels $t < t'$,

$$\begin{aligned} & \mathbb{E} \|s_t(X_t \mid t', X_{t'}) - s_t(X_t)\|^2 \\ &= \frac{1}{t^2} (\text{MMSE}(t') - \text{MMSE}(t)), \end{aligned} \quad (65)$$

with the sign understood according to the direction in which the diffusion trajectory is traversed. Thus, for a sufficiently fine discretisation with step size

$$\Delta t_k := t_{k+1} - t_k, \quad (66)$$

the contribution of the k -th interval is, up to higher-order terms,

$$\mathcal{E}_k \asymp \frac{|\text{MMSE}'(t_k)|}{t_k^2} (\Delta t_k)^2. \quad (67)$$

Consequently, the total discretisation error is approximated by

$$\mathcal{E}_{\text{disc}} \asymp \sum_{k=0}^{K-1} \alpha(t_k) (\Delta t_k)^2, \quad \alpha(t) := \frac{|\text{MMSE}'(t)|}{t^2}. \quad (68)$$

Optimal allocation of time steps. We now choose the discretisation points to minimise (68) subject to a fixed number K of steps. Introduce the local step density

$$n(t) := \frac{1}{\Delta t(t)}. \quad (69)$$

In the continuum limit,

$$K \approx \int n(t) dt, \quad (70)$$

while

$$\begin{aligned} \mathcal{E}_{\text{disc}} &\approx \int \alpha(t) \Delta t(t) dt \\ &= \int \frac{\alpha(t)}{n(t)} dt. \end{aligned} \quad (71)$$

Minimising this functional subject to $\int n(t) dt = K$ gives

$$n(t) \propto \sqrt{\alpha(t)}, \quad (72)$$

and therefore

$$\Delta t(t) \propto \frac{1}{\sqrt{\alpha(t)}} = \frac{t}{\sqrt{|\text{MMSE}'(t)|}}. \quad (73)$$

Finally, using Equation 57,

$$\text{MMSE}'(t) = \mathbb{E} \|A_t(Y_t)\|_F^2, \quad (74)$$

so the resulting MMSE-adaptive discretisation is

$$\Delta t(t) \propto \frac{t}{\sqrt{\mathbb{E} \|A_t(Y_t)\|_F^2}}. \quad (75)$$

6.4 Interpretation

Geometric Interpretation. We see the effect of the $\text{MMSE}'(t)$ term on time-step density assignment. Notably, since $t d_{\text{eff}}(t) = \text{MMSE}(t)$, we have $\text{MMSE}'(t) = d_{\text{eff}}(t) + t d'_{\text{eff}}(t)$. Hence for higher d_{eff} and crucially changing d_{eff} for larger time we require smaller time steps. Furthermore, $\mathbb{E} \|A_t(Y_t)\|_F^2 = \sum_i \mathbb{E} [a_i(t)^2]$, so this discretisation assigns a high density of small timesteps whenever we have expansive eigenvalues $a_i > 1$, in expectation.




Figure 3: Functionals and assigned time step evolution for Bernouli($\{+1, -1\}$), Uniform($\mathcal{S}^1$), and Uniform($\{\mathcal{S}^1\} \times \{\frac{1}{4}, -\frac{1}{4}\}$)

Information Theortic Interpretation. By (Guo et al., 2005) we have

$$\frac{d}{dt}I(X_0; X_t) = \frac{dI}{d\gamma} \cdot \frac{d\gamma}{dt} = \left(\frac{1}{2} \text{MMSE}(t)\right) \left(-\frac{2}{t^3}\right) = -\frac{\text{MMSE}(t)}{t^3}$$

where γ is the signal to noise ratio (here $1/t^2$). More to be said!

Probabilistic Interpretation. To be discussed.

Acknowledgements

I am deeply grateful to my supervisors Eddie Aamari and Simon Coste, who are the sole reason this research internship was possible and helped me explore a question I had for more than a year.

AI Usage

Image to text converters were used to aid in typesetting equations, and large language models were used for some code generation and in editing parts of the report.

References

- Eddie Aamari and Clément Berenfeld. A theory of stratification learning. *arXiv preprint arXiv:2405.20066*, 2024.
- Michael Albergo, Nicholas M Boffi, and Eric Vanden-Eijnden. Stochastic interpolants: A unifying framework for flows and diffusions. *Journal of Machine Learning Research*, 26(209):1–80, 2025.
- Joe Benton, Valentin De Bortoli, Arnaud Doucet, and George Deligiannidis. Nearly d -linear convergence bounds for diffusion models via stochastic localization. In *International Conference on Learning Representations*, volume 2024, pages 36916–36936, 2024.
- Bradley CA Brown, Anthony L Caterini, Brendan Leigh Ross, Jesse C Cresswell, and Gabriel Loaiza-Ganem. The union of manifolds hypothesis. In *NeurIPS 2022 Workshop on Symmetry and Geometry in Neural Representations*, 2022a.
- Bradley CA Brown, Anthony L Caterini, Brendan Leigh Ross, Jesse C Cresswell, and Gabriel Loaiza-Ganem. Verifying the union of manifolds hypothesis for image data. *arXiv preprint arXiv:2207.02862*, 2022b.
- Herbert Federer. Curvature measures. *Transactions of the American mathematical Society*, 93(3):418–491, 1959.
- Christopher R Genovese, Marco Perone-Pacifco, Isabella Verdinelli, and Larry Wasserman. Nonparametric ridge estimation. 2014.
- Dongning Guo, Shlomo Shamai, and Sergio Verdú. Mutual information and minimum mean-square error in gaussian channels. *IEEE transactions on information theory*, 51(4):1261–1282, 2005.
- Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in neural information processing systems*, 33:6840–6851, 2020.
- Jiadong Liang, Zhihan Huang, and Yuxin Chen. Low-dimensional adaptation of diffusion models: Convergence in total variation. *arXiv preprint arXiv:2501.12982*, 2025.
- Yaron Lipman, Ricky TQ Chen, Heli Ben-Hamu, Maximilian Nickel, and Matt Le. Flow matching for generative modeling. *arXiv preprint arXiv:2210.02747*, 2022.
- Peter Potapchik, Iskander Azangulov, and George Deligiannidis. Linear convergence of diffusion models under the manifold hypothesis. *arXiv preprint arXiv:2410.09046*, 2024.
- Jiaming Song, Chenlin Meng, and Stefano Ermon. Denoising diffusion implicit models. *arXiv preprint arXiv:2010.02502*, 2020a.

- Yang Song and Stefano Ermon. Generative modeling by estimating gradients of the data distribution. *Advances in neural information processing systems*, 32, 2019.
- Yang Song, Jascha Sohl-Dickstein, Diederik P Kingma, Abhishek Kumar, Stefano Ermon, and Ben Poole. Score-based generative modeling through stochastic differential equations. *arXiv preprint arXiv:2011.13456*, 2020b.

Appendix: MMSE Evolution

Theorem 6.9. Let $X \in \mathbb{R}^D$ be a random vector with probability density $p(x)$ and finite second moment, and let $Z \sim \mathcal{N}(0, I_D)$ be independent of X . Define the observation process $Y_t = X + \sqrt{t}Z$ for $t > 0$, and the Minimum Mean Squared Error function:

$$\text{MMSE}(t) = \mathbb{E} [\|X - \mathbb{E}[X | Y_t]\|^2]$$

Then $\text{MMSE}(t)$ is differentiable with respect to $t > 0$, and its derivative satisfies:

$$\text{MMSE}'(t) = \frac{1}{t^2} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2)]$$

Proof. The marginal density of Y_t , $p_t(y) = \int p(x)(2\pi t)^{-D/2} \exp\left(-\frac{\|y-x\|^2}{2t}\right) dx$, satisfies the heat equation $\frac{\partial p_t(y)}{\partial t} = \frac{1}{2} \Delta_y p_t(y)$. Define the score function $s_t(y) = \nabla_y \log p_t(y)$ and the Fisher Information trace $J(t) = \mathbb{E} [\|s_t(Y_t)\|^2]$.

1. Score representation of conditional statistics. Spatial differentiation of $p_t(y)$ yields Tweedie's formula:

$$s_t(y) = \frac{\nabla_y p_t(y)}{p_t(y)} = -\frac{y - \mathbb{E}[X | Y_t = y]}{t} \implies X - \mathbb{E}[X | Y_t] = -ts_t(Y_t)$$

Taking the spatial Jacobian of $\mathbb{E}[X | Y_t = y] = y + ts_t(y)$:

$$\nabla_y \mathbb{E}[X | Y_t = y] = \frac{1}{t} \text{Cov}(X | Y_t = y) = I_D + t \nabla_y s_t(y)$$

Rearranging yields the score Hessian:

$$\nabla_y s_t(y) = \frac{\text{Cov}(X | Y_t = y) - tI_D}{t^2}$$

2. Relation between MMSE and Fisher Information. By Tweedie's formula:

$$\text{MMSE}(t) = \mathbb{E} [\| -ts_t(Y_t) \|^2] = t^2 J(t)$$

Applying integration by parts ($\mathbb{E}[\|s_t(Y_t)\|^2] = -\mathbb{E}[\text{tr}(\nabla_y s_t(Y_t))]$) to the score Hessian:

$$\begin{aligned} J(t) &= -\mathbb{E} \left[\text{tr} \left(\frac{\text{Cov}(X | Y_t) - tI_D}{t^2} \right) \right] = \frac{D}{t} - \frac{\text{MMSE}(t)}{t^2} \\ \implies \text{MMSE}(t) &= tD - t^2 J(t) \end{aligned}$$

3. Time evolution of Fisher Information. Differentiating $J(t) = \int p_t(y) \|\nabla \log p_t(y)\|^2 dy$ with respect to t under heat flow $\frac{\partial p_t}{\partial t} = \frac{1}{2} \Delta p_t$ yields de Bruijn's score identity:

$$J'(t) = -\mathbb{E} [\|\nabla_y s_t(Y_t)\|_F^2] = -\mathbb{E} [\text{tr} ((\nabla_y s_t(Y_t))^2)]$$

Substituting the expression for $\nabla_y s_t(Y_t)$:

$$\begin{aligned} J'(t) &= -\mathbb{E} \left[\text{tr} \left(\left(\frac{\text{Cov}(X | Y_t) - t I_D}{t^2} \right)^2 \right) \right] \\ &= -\frac{1}{t^4} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2 - 2t \text{Cov}(X | Y_t) + t^2 I_D)] \\ &= -\frac{1}{t^4} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2)] + \frac{2}{t^3} \text{MMSE}(t) - \frac{D}{t^2} \end{aligned}$$

4. Total differentiation of MMSE(t). Differentiating $\text{MMSE}(t) = tD - t^2 J(t)$ via the product rule:

$$\text{MMSE}'(t) = D - 2tJ(t) - t^2 J'(t)$$

Substituting $2tJ(t) = 2D - \frac{2\text{MMSE}(t)}{t}$ and $-t^2 J'(t)$:

$$\begin{aligned} \text{MMSE}'(t) &= D - \left(2D - \frac{2\text{MMSE}(t)}{t} \right) + \left(\frac{1}{t^2} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2)] - \frac{2\text{MMSE}(t)}{t} + D \right) \\ &= (D - 2D + D) + \left(\frac{2\text{MMSE}(t)}{t} - \frac{2\text{MMSE}(t)}{t} \right) + \frac{1}{t^2} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2)] \\ &= \frac{1}{t^2} \mathbb{E} [\text{tr} (\text{Cov}(X | Y_t)^2)] \end{aligned}$$

□